

DEEP LEARNING FOR GRAPHS AND SETS

Adriana Romero

March 26th, 2019 @ UPC School



Bachelor



UNIVERSITAT
BARCELONA

MSc



UNIVERSITAT
RESEARCH GROUP



UNIVERSITAT
BARCELONA

PhD




McGill

Adjunct Professor



Research Scientist



MILA
Université
de Montréal

post-doc



Model Compression (FitNets)

Unsupervised pre-training (EPLS)

- Conditional generation
- Data multi-modality
- Graphs and sets


- Biomedical challenges:
- High dimensional data (genomics)
 - Structured prediction (imaging)
 - Graph-structured data (meshes)

- Applications:
- Image classification
 - Semantic segmentation
 - Remote sensing

PREREQUISITES

- MLP
- CNN
- RNN
- Attention
- How to train neural networks

DEEP LEARNING MODELS & APPLICATIONS



OUTLINE

GRAPH-STRUCTURED DATA:

- Motivation and problem formulation
- Overview of prior work
- Graph attention networks
- Results
- Other graph applications

SET PREDICTION

- Motivation and problem formulation
- Deep learning-based set prediction models, losses and cardinality estimation
- Results

RECIPE GENERATION

- Motivation and problem formulation
- Model
- Results

GRAPH-STRUCTURED DATA

P. Velickovic, G. Cucurull, A. Casanova, A. Romero, P. Lio, Y. Bengio


Graph Attention Networks @ ICLR 2018




UNIVERSITY OF
CAMBRIDGE



MOTIVATION




Eric Wipbski



Graphs are everywhere!

PROBLEM FORMULATION



Transductive learning:


Node classification problem:

Input: a matrix of node features F and an adjacency matrix A

Output: a matrix of node class probabilities


We assume for simplicity, that edges are **unweighted** and **undirected**.

PROBLEM FORMULATION



Transductive learning: training sees the features of *all* nodes.

PROBLEM FORMULATION



Node classification problem:

Input: a matrix of node features F and an adjacency matrix A
Output: a matrix of node class probabilities

We assume for simplicity, that edges are unweighted and undirected.

Transductive learning: training sees the features of *all* nodes.

Inductive learning: training does *not* have access to all nodes upfront.

NEURAL NETWORKS FOR GRAPHS

per-node classifier

Does not exploit graph structure!

NEURAL NETWORKS FOR GRAPHS

per-node classifier

per-node classifier +
node similarity constraint
(Weston et al., 2008)

per-node classifier
based on node & structure features
(Perozzi et al., 2014; Tang et al., 2015;
Grover et al., 2016; Yang et al., 2016)

NEURAL NETWORKS FOR GRAPHS

per-node classifier

per-node classifier +
node similarity constraint
(Weston et al., 2008)

Graph structure injected indirectly!

per-node classifier
based on node & structural features
(Perozzi et al., 2014; Tang et al., 2015;
Grover et al., 2016; Yang et al., 2016)



A good representation of a node should allow us to easily predict the nodes that *surround* it.

NEURAL NETWORKS FOR GRAPHS

per-node classifier

per-node classifier +
node similarity constraint
(Weston et al., 2008)

per-node classifier
based on node & structural features
(Perozzi et al., 2014; Tang et al., 2015;
Grover et al., 2016; Yang et al., 2016)

Graph Neural Networks

(Gori et al., 2005; Scarselli et al., 2009)

Gated Graph Neural Networks
(Li et al., 2016)

Our graphs have static features!




A good representation of a node should allow us to easily predict the nodes that *surround* it.

HOW ABOUT LEVERAGING CONVOLUTIONS? (1)

CNN work well on data defined in n-D grids



2D grid



“I like cats and dogs.”

1D grid



HOW ABOUT LEVERAGING CONVOLUTIONS? (2)

Why convolutions?

Independent of input size


Fixed # parameters (weight sharing)

Localized (act on local neighborhoods)

Specify different importance per node

Applicable to both transductive and

inductive problems



HOW ABOUT LEVERAGING CONVOLUTIONS? (2)

Why convolutions?

Independent of input size

Fixed # parameters (weight sharing)

Localized (act on local neighborhoods)

Specify different importance per node

Applicable to both transductive and

inductive problems

The diagram illustrates a convolution operation. On the left is a 7x7 input matrix with values ranging from 0 to 4. A 3x3 kernel is shown in the center, highlighted with a blue border. Dotted lines connect the kernel to its corresponding receptive field in the input matrix. The result of the convolution is shown on the right, where the output values are 1, 2, 3, 4, and 5. The result is highlighted with a green border.

0	1	1	1	0	0	0
0	0	1	1	1	0	0
0	0	0	1	1	1	0
0	0	0	1	1	0	0
0	0	1	1	0	0	0
0	1	1	0	0	0	0
1	1	0	0	0	0	0

*

1	0	1
0	1	0
1	0	1

=

1	4	3	4	1
1	2	4	3	3
1	2	3	4	1
1	3	3	1	1
3	3	1	1	0

HOW ABOUT LEVERAGING CONVOLUTIONS? (3)

Why convolutions?

Independent of input size

Fixed # parameters (weight sharing)


Localized (act on local neighborhoods)

Specify different importance per node

Applicable to both transductive and

inductive problems

How do we apply a convolutional filter over a graph?




?

w^0	w^1	w^2
w^3	w^4	w^5
w^6	w^7	w^8

GRAPH SPECTRAL NETWORKS

Operating in the graph **spectral** domain



$$\mathbf{L} = \begin{bmatrix} 2 & -1 & 0 & 0 & -1 & 0 \\ -1 & 3 & -1 & 0 & -1 & 0 \\ 0 & -1 & 2 & -1 & 0 & 0 \\ 0 & 0 & -1 & 3 & -1 & -1 \\ -1 & -1 & 0 & -1 & 3 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 \end{bmatrix}$$

Graph Laplacian:

D (degree matrix) – A (adjacency matrix)
(or a *normalized* version of it)

$$\mathbf{L} = \mathbf{U} \Lambda \mathbf{U}^T$$

- Multiplying the feature matrix by \mathbf{U}^T allows us to enter the spectral domain of the graph.
- Convolution amounts to simple multiplication.

$$\vec{h}'_i = \mathbf{U} \left(\vec{w} \odot \mathbf{U}^T \mathbf{w} \vec{h}_i \right)$$

- Computing \mathbf{U} can be expensive.
- Filters are not localized.

(Bruna et al., 2014)

CHEBYNETS & GCN

Rather than computing the Fourier transform, use Chebyshev polynomials of order k to approximate it.

$$\vec{h}'_i = \sum_{k=0}^K w_k T_k(\mathbf{L}) \mathbf{W} \vec{h}_i$$

GCN set to $K = 1$, further simplifying the definition of a convolutional layer.

- These polynomials have a recursive formulation, highly simplifying the computation.
 - T_k is a weighted sum of all powers of L up to L^k . \longrightarrow K-localized
 - The number of parameters is fixed (k).
- **Unable to specify different weights to different nodes in a neighborhood.**

(Defferrard et al., 2016; Kipf & Welling, 2017)

NON-SPECTRAL APPROACHES

✓ **Do not assume a particular & fixed graph Laplacian**

Molecular fingerprint networks

GraphSAGE


- Processing with various degrees by learning a **separate** weight matrix **per node degree**.
- **Does not scale to graphs with very wide degree distributions.**
- Restricts every degree to be the same, by sampling a **fixed-size** of neighbors.
- **Inherently drops relevant data.**

(Duvenaud et al., 2015; Hamilton et al., 2017)

**ATTENTION AS A WAY TO OVERCOME THE
PREVIOUS ISSUES**


ATTENTION (1)

Attention allows us to focus on a subset of information and select the most pertinent piece of it.



Attention input, can be of variable length
(image features, seq of annotations)

ATTENTION (2)



1. Combine input and hidden information
$$\mathbf{m}_i = f(\mathbf{W}_c \mathbf{c} + \mathbf{W}_y \mathbf{y}_i)$$
2. Apply softmax to obtain attention coefficients s_i
3. Apply attention coefficients to input and aggregate

$$\mathbf{z} = \sum_i s_i \odot \mathbf{y}_i$$

GRAPH ATTENTION NETWORKS (1)


The idea is to leverage the **self-attention** operator to *emulate* convolutions on graphs.

Self attention allows the input to attend **over itself**:

$$\begin{aligned}\alpha_{ij} &= a(\vec{h}_i, \vec{h}_j) \\ \vec{h}'_i &= \sum_j softmax_j(\alpha_{ij}) \vec{h}_j\end{aligned}$$

- A naïve formulation would compute attention coefficients over all pairs of nodes, dropping all structural information.
- We restrict the model to only attend over the node's neighborhood, making the operator localized.


GRAPH ATTENTION NETWORKS (2)



$$e_{ij} = \vec{a}(\vec{W}\vec{h}_i, \vec{W}\vec{h}_j)$$

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k \in \mathcal{N}_i} \exp(e_{ik})}$$

GRAPH ATTENTION NETWORKS (3)



$$\vec{h}'_i = \sigma \left(\sum_{j \in \mathcal{N}_i} \alpha_{ij} \mathbf{W} \vec{h}_j \right)$$


For one attention head.

DATASETS (1)


	Cora	Citeseer	Pubmed
# Graphs	1	1	1
# Nodes	2708	3327	19717
# Edges	5429	4732	44338
# Features / Node	1433	3703	500
# Classes	7	6	3
# Training Nodes	140	120	60
# Validation Nodes	500	500	500
# Test Nodes	1000	1000	1000

Nodes: documents Features: sparse BoW
Edges: citations Output: 1 class per document


RESULTS (1)



RESULTS (1)



RESULTS (1)



DATASETS (2)

	PPI	HCP
# Graphs	24	100
# Nodes	56944	119500
# Edges	818716	342900
# Features / Node	50	9
# Classes	121 (multilabel)	3
# Training Nodes	44906 (20 graphs)	95600 (80 graphs)
# Validation Nodes	6514 (2 graphs)	11950 (10 graphs)
# Test Nodes	5524 (2 graphs)	11950 (10 graphs)

Graphs: one per human tissue

Features: positional gene sets, motif gene sets, immunological signatures

Output: n active labels (protein roles)

DATASETS (2)


	PPI	HCP
# Graphs	24	100
# Nodes	56944	119500
# Edges	818716	342900
# Features / Node	50	9
# Classes	121 (multilabel)	3
# Training Nodes	44906 (20 graphs)	95600 (80 graphs)
# Validation Nodes	6514 (2 graphs)	11950 (10 graphs)
# Test Nodes	5524 (2 graphs)	11950 (10 graphs)

Meshes: one per subject


Features: structural features, functional features

Output: area 44, area 45, none

RESULTS PPI



RESULTS CORTICAL MESHES



(Cucurull et al., 2018)

GRAPHS WRAP UP

- We introduced GAT, a model which:
 - Implicitly allows **assigning different weights** to nodes in a same neighborhood.
 - Provides a trivially **localized** operator.
 - Does not depend on upfront access to the graph structure (**inductive problems**).
 - Has a **fixed number of parameter** (dependent on the feature count, not on node/degree counts).
 - **Computationally efficient:** self-attention parallelizable across edges, computation of output features parallelizable across nodes.
- We highlighted its potential in 3 citation network datasets as well as 2 applications in the biomedical domain.

OTHER GRAPH APPLICATIONS


3D UNDERSTANDING




(a) Voxels
(262,144 units)



(b) Point cloud
(30,000 points)



(c) Uniform mesh
(2416 vertices)




(d) Adaptive mesh
(120 vertices)


Figure credit: Edward Smith

3D UNDERSTANDING TASK

3D reconstruction from a single view image.




UNIFORM MESH RECONSTRUCTION




Trained to minimize the Chamfer loss:

$$\sum_{p \in S} \min_{q \in \hat{S}} \|p - q\|_2^2 + \sum_{p \in \hat{S}} \min_{q \in S} \|p - q\|_2^2$$

Figure credit: Edward Smith



ADAPTIVE MESH RECONSTRUCTION




Trained with:

- Point-to-surface loss, accounting for the surface defined by the vertices, rather than their position.
- Global mesh loss, comparing embeddings of the ground truth mesh vs the predicted one.

(Smith et al., 2019)

ADAPTIVE MESH RESULTS (1)



ADAPTIVE MESH RESULTS (2)



Input
image



(Smith et
al., 2019)



(Wang et
al., 2018)



Input
image




(Smith et
al., 2019)



(Wang et
al., 2018)

ADAPTIVE MESH RESULTS (3)



SET PREDICTION

L. Pineda*, A. Salvador*, M. Drozdzal, A. Romero



MOTIVATION (1)

Among image understanding tasks, image classification has arguably received the most attention.



http://image-net.org/challenges/talks_2017/ILSVRC2017_overview.pdf

MOTIVATION (2)

However... everyday life pictures are typically complex scenes, with **multiple labels** per image.



dog,
person,
chair, etc



bottle,
chair,
person,
table, etc



sky, tree,
building,
car, tree,
fence,
bench, etc




clouds,
sky,
sunset



ham,
tomato,
bagel,
vinaigrette

MULTI-LABEL CLASSIFICATION



Framed as an **image-to-set** (of labels) prediction problem:

- Image labels may exhibit relevant **dependencies** (co-occurrences)
- The number of labels per image is **variable**
- **Order** of labels should not matter

DL SET PREDICTION MODELS (1)






Image feature extractor


DL SET PREDICTION MODELS (2)




DL SET PREDICTION MODELS (3)



DL SET PREDICTION MODELS (4)



MLP: PREDICTION LOSSES (1)



If we consider **independent** outputs:

- ✓ \hat{s} represents a vector of independent Bernoulli distributions (sigmoid non-linearity)
- ✓ We could use binary cross-entropy as loss function:


$$-\frac{1}{L} \sum_l s_l \log \hat{s}_l + (1 - s_l) \log(1 - \hat{s}_l),$$

where s_l is the ground truth.

- ✓ We could use a structured loss (soft IoU):

$$\frac{\sum_l s_l \hat{s}_l}{\sum_l s_l + \hat{s}_l - s_l \hat{s}_l}$$

MLP: PREDICTION LOSSES (2)




If we consider **dependent** outputs:

- ✓ \hat{s} represents a categorical distribution (softmax non-linearity)
- ✓ What could we use as a target?
If the ground truth is [1, 0, 0, 1, 0], we could use [0.5, 0, 0, 0.5, 0] as s
- ✓ Then, we could use cross-entropy as proxy loss function:

$$-\frac{1}{L} \sum_l s_l \log \hat{s}_l$$

MLP: CARDINALITY PREDICTION



- ✓ We could use \hat{s} to determine the cardinality of the set:


$$\hat{s}_l > 0.5 \forall l$$

- ✓ We could predict the cardinality from a second output:

- MSE (ReLU non-linearity)

- Categorical cross-entropy (softmax non-linearity)

AUTO-REGRESSIVES: PREDICTION LOSSES (1)




If we consider **ordered** outputs:

- ✓ \hat{s}_k (at time step k) represents a categorical distribution (softmax non-linearity)
- ✓ We could use categorial cross-entropy as loss function for each time step:

$$-\frac{1}{L} \sum_l s_k \log \hat{s}_k ,$$


where s_k is a one hot vector containing one of the ground truth labels.

AUTO-REGRESSIVES: PREDICTION LOSSES (2)



If we consider no order:

- ✓ Pool across time steps:




- ✓ We could use binary cross-entropy as loss function

$$-\frac{1}{L} \sum_l s_l \log \hat{s}_l + (1 - s_l) \log(1 - \hat{s}_l),$$

where s_l is the ground truth

AUTO-REGRESSIVES: CARDINALITY PREDICTION



If we consider **ordered** outputs:

- ✓ end-of-sequence (*eos*) token predicted as last step
- ✓ *eos* step included in categorical cross-entropy loss

If we consider **no ordered**:

- ✓ *eos* token predicted as last step
- ✓ Additional loss at each time step to decide whether *eos* should be predicted (categorical cross-entropy)

HOW DO ALL THESE METHODS COMPARE?

DATASETS



VOC 2007

20 categories
4.9k train images
512 valid images
4.9k test images
All (partially)
visible objects

MS COCO 2014

80 categories
74.5k train images
8.3k valid images
40.5k test images
All (partially)
visible objects

ADE20k

150 categories
18.2k train images
2k valid images
2k test images
All (partially)
visible objects


NUS-WIDE

81 categories
45.6k train images
16.2k valid images
107.9k test images
54.5k test images


Recipe1M

1486 categories
252.5k train images
5k valid images
54.5k test images

RESULTS (1)



RESULTS (2)



SETS WRAP UP

- We presented a **comprehensive analysis** of methods suitable for image-to-set prediction:
 - We evaluated their performance in **5 diverse datasets**
 - Using a **uniform set of metrics**
 - **Comparable budgets** for hyperparameter tuning
- Our analysis suggests that **auto-regressive models are better choices** than feed-forward models for the task, performing consistently well across all considered datasets.
 - They inherently handle **set cardinality prediction, label co-occurrences**.
 - **Shuffling** label order tends to increase performance.

RECIPE GENERATION

A. Salvador, M. Drozdzal, X. Giro-Nieto, A. Romero

Inverse Cooking: Recipe Generation from Food Images @ CVPR 2019



MOTIVATION




Shapes and colors

Invisible ingredients



More abstract understanding of what food is.

WHY GENERATING RECIPES?




Previous work focuses on retrieving recipes for a query image. But...

- Retrieval is **constrained** to database recipes (inaccurate matches).
- Framed to retrieve ingredients and recipes **as a whole**.
- Prohibits **user manipulation** (e.g. ingredient replacement).

PROBLEM FORMULATION


Generate ingredients list & cooking instructions.



Ingredient's list treated as **sets**.

- Order **invariance**.
- **Variable** cardinality.
- **Dependencies** among different elements in the set.

RECIPE GENERATION FROM IMAGE AND INGREDIENTS' SET



RECIPE GENERATION FROM IMAGE AND INGREDIENTS' SET







Figure credit: Amaia Salvador

DATASET: RECIPE 1M

A structured dataset of 1M recipes and 800k images scraped from ~20 cooking recipe websites.



The average cooking recipe in our dataset contains 9 ingredients which are transformed over the course of 10 instructions.

Salvador, Hynes et al., 2017

QUALITATIVE RESULTS (1)



beef, onion, tomato,
water, pasta, beans,
cheese, chili



sugar, flour, salt,
baking_soda, butter,
egg, extract



edamame, pepper, beans,
corn, dill, juice,
italian_dressing, avocado

tomato, onion, beans,
pepper, broth, clove,
carrot, oil, salt, chili

egg, flour, sugar, butter,
salt, cinnamon,
baking_powder, milk

beans, corn, pepper, tomato,
onion, vinegar, oil

Figure credit: Amaia Salvador

QUALITATIVE RESULTS (2)



Title: Basic sugar cookies

Ingredients

- flour
- egg
- butter
- salt
- sugar

Instructions

1. Preheat oven to 350.
2. In a mixing bowl **cream butter** with **sugar**.
3. Add the **eggs** one at a time.
4. Then add the **flour** and **salt** and mix thoroughly.
5. Drop by teaspoonfuls on a baking sheet and bake for 8-10 minutes (or until golden brown).
6. Remove from oven and transfer to a wire rack to cool.

QUALITATIVE RESULTS (3)



Title: Macaroni and cheese with beans

Ingredients

- onion
- pepper
- cheese
- salt
- beans
- clove
- pasta
- bacon
- tomato
- oil
- water

Instructions

1. Cook pasta according to package directions.
2. Meanwhile, heat olive oil in a deep skillet over medium high heat.
3. Saute onion and garlic in olive oil until tender, 7 to 10 minutes.
4. Stir in the beans, bacon, salt and pepper.
5. Saute, stirring, until beans are heated through, 10 to 12 minutes.
6. Stir in water and tomato sauce, bring to a boil, then reduce heat and simmer, covered, until sauce is thickened, about 10 minutes.
7. Stir in cheese.
8. Drain pasta and place on serving plates.

RECIPE GENERATION WRAP UP

We coupled the set prediction module with a conditional transformer, which predicts a recipe from an input image and the predicted ingredients' set, showcasing compelling results.

COLLABORATORS



Guillem Cururull
Arantxa Casanova



Amaia Salvador
Xavier Giró-Nieto



Konrad Wagstyl
Estrid Jakobsen
Alan Evans



Michal Drozdzal
Luis Pineda



Edward Smith
Scott Fujimoto
Dave Meger



UNIVERSITY
OF
CAMBRIDGE

Petar Velickovic
Pietro Lio
Konrad Wagstyl

DEEP LEARNING FOR GRAPHS AND SETS

Adriana Romero

March 26th, 2019 @ UPC School

THANKS!