

# Machine Translation

Cristina España-Bonet

UdS & DFKI, Saarbrücken, Germany

Artificial Intelligence with Deep Learning

2nd May 2019

## Neural Machine Translation reaches historic milestone: human parity for Chinese to English translations

Rate this article ★★★★☆

 Microsoft Translator March 14, 2018

 Share 12  7  0  0



# Achieving Human Parity on Automatic Chinese to English News Translation

Hany Hassan, Anthony Aue, Chang Chen, Vishal Chowdhary, Jonathan Clark,  
Christian Federmann, Xuedong Huang, Marcin Junczys-Dowmunt, William Lewis,  
Mu Li, Shujie Liu, Tie-Yan Liu, Renqian Luo, Arul Menezes, Tao Qin,  
Frank Seide, Xu Tan, Fei Tian, Lijun Wu, Shuangzhi Wu, Yingce Xia,  
Dongdong Zhang, Zhirui Zhang, and Ming Zhou

Microsoft AI & Research

## Abstract

Machine translation has made rapid advances in recent years. Millions of people are using it today in online translation systems and mobile applications in order to communicate across language barriers. The question naturally arises whether such systems can approach or achieve parity with human translations. In this paper, we first address the problem of how to define and accurately measure human parity in translation. We then describe Microsoft’s machine translation system and measure the quality of its translations on the widely used WMT 2017 news translation task from Chinese to English. We find that our latest neural machine translation system has reached a new state-of-the-art, and that the translation quality is at human parity when compared to professional human translations. We also find that it significantly exceeds the quality of crowd-sourced non-professional translations.

## 1 Introduction

Recent years have seen human performance levels reached or surpassed in tasks ranging from games such as Go [32] to classification of images in ImageNet [20] to conversational speech recognition on the Switchboard task [49].

In the area of machine translation, we have seen dramatic improvements in quality with the advent of attentional encoder-decoder neural networks [34, 3, 38]. However, translation quality continues to vary a great deal across language pairs, domains, and genres, more or less in direct

# Motivation

## Applications

SOCIETAT > **CASTELLERS** CIÈNCIA EDUCACIÓ MEDI AMBIENT TEMPS SANITAT SUCCE



### Tot a punt per al concurs de castells de Tarragona

CASTELLERS



Els Castellers de Viladecans celebren el seu 5è aniversari

ENSURT SENSE CONSEQUÈNCIES



Una enxaneta se salva miraculosament d'una caiguda a la plaça de Sant Jaume

CANVI DE JUNTA

Minyons demana girar full i treballar en cooperació amb la Coordinadora de Colles Castelleres

La colla egarenca advoca per iniciar una "nova etapa" arcadà per la "leialtat reciproca i...

DIADA CASTELLERA



Minyons, Capgrossos i Borinots han de renunciar les seves màximes fites en

SOCIEDAD > **CASTELLERS** CIENCIA EDUCACIÓN MEDIO AMBIENTE TIEMPO SANIDAD SUCE



### Todo a punto para el concurso de castells de Tarragona

REIVINDICACIÓN



Los Castellers de Barcelona descargan una torre solo de mujeres

CASTELLERS



Los Castellers de Viladecans celebran su 5º aniversario

SUSTO EN LA PLAZA DE SANT JAUME



Una 'enxaneta' se salva tras una caída en Barcelona

CAMBIO DE JUNTA

Minyons pide pasar página y trabajar en cooperación con la Coordinadora de Colles Castelleres

La colla egarenca advoca por iniciar una "nueva etapa" arcadà per la "leialtat reciproca y..."

# Outline

- 1 Introduction
- 2 Statistical Machine Translation
- 3 Neural Machine Translation
- 4 Software & References

# Introduction

1 Introduction

2 Statistical Machine Translation

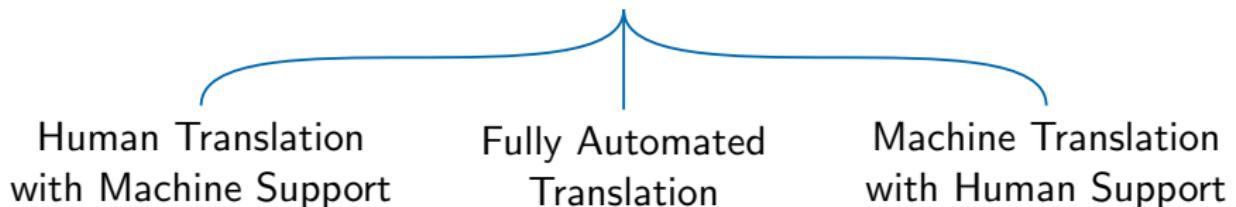
3 Neural Machine Translation

4 Software & References

# Introduction

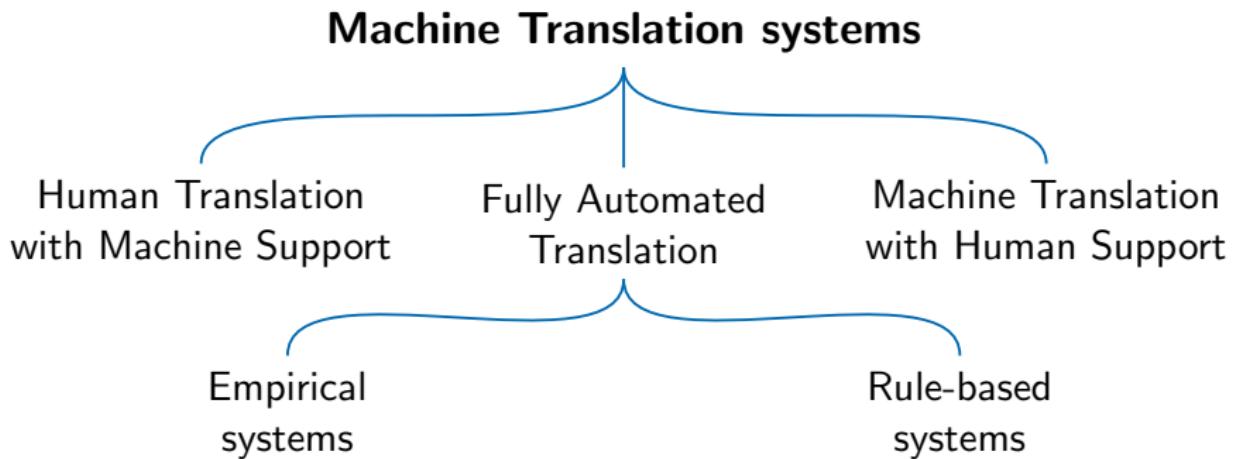
## *Naïve Taxonomy*

### **Machine Translation systems**



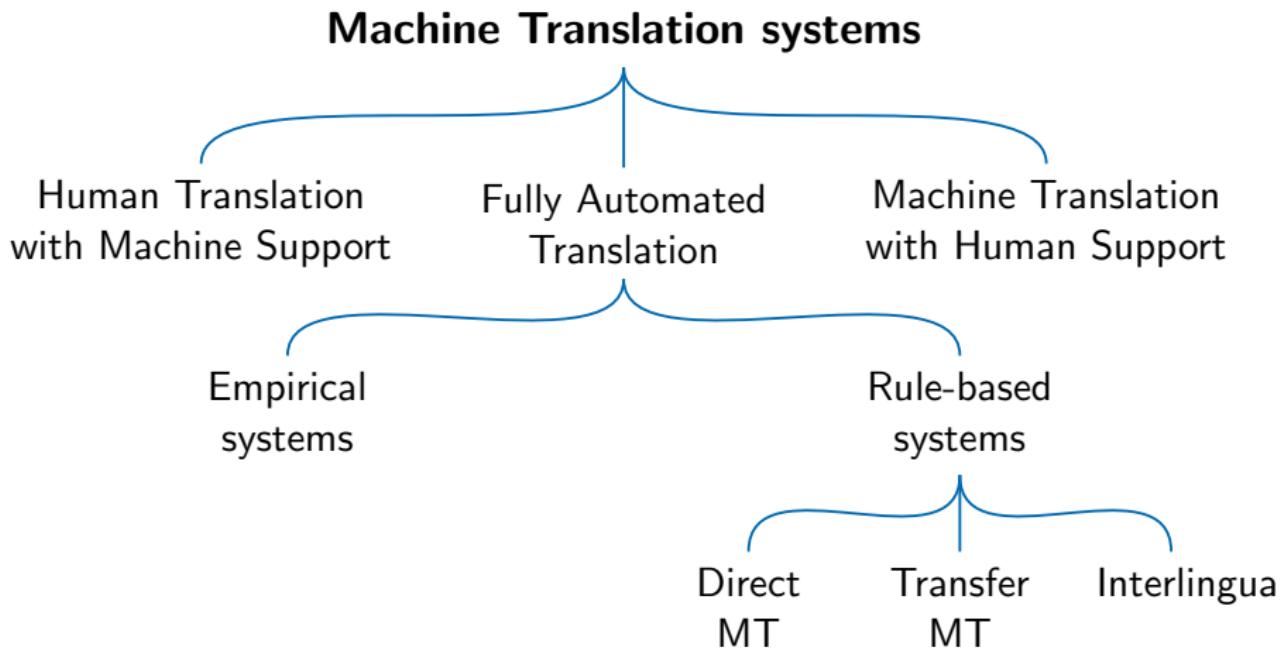
# Introduction

## *Naïve Taxonomy*



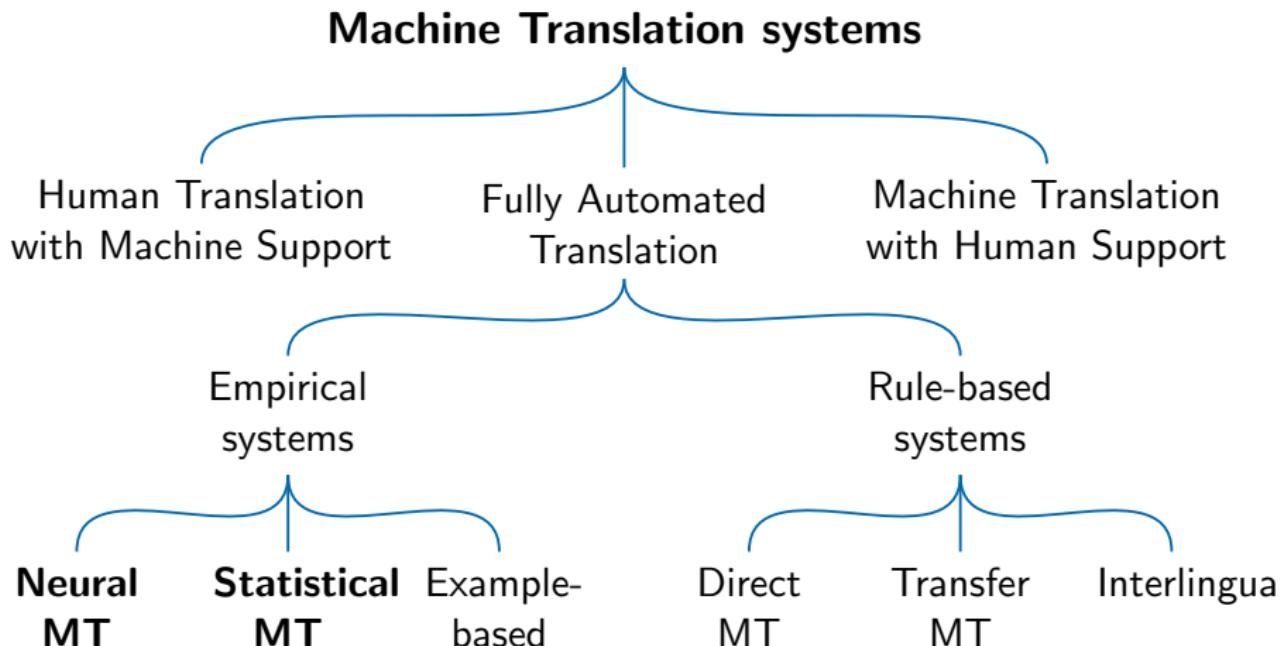
# Introduction

## *Naïve Taxonomy*



# Introduction

## *Naïve Taxonomy*

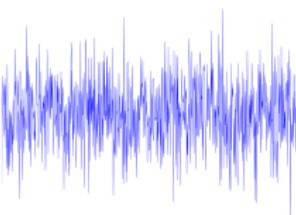


# Statistical Machine Translation

## *The Noisy Channel*

**The Noisy Channel** as a statistical approach to translation:

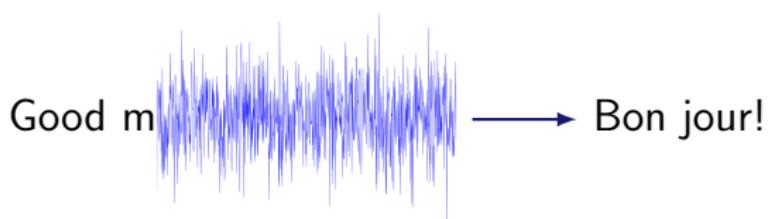
Good morning! →



# Statistical Machine Translation

## *The Noisy Channel*

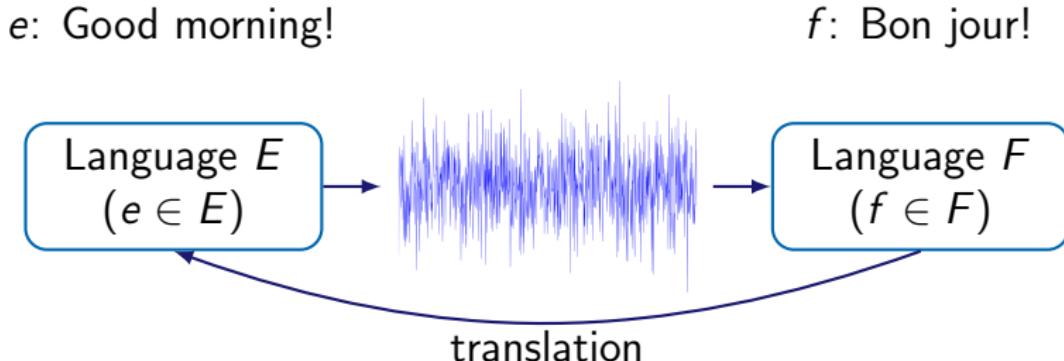
**The Noisy Channel** as a statistical approach to translation:



# Statistical Machine Translation

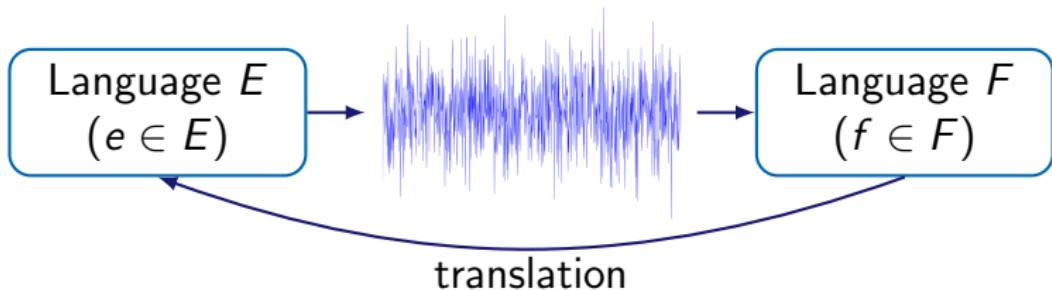
## The Noisy Channel

**The Noisy Channel** as a statistical approach to translation:



# Statistical Machine Translation

## *The Noisy Channel*

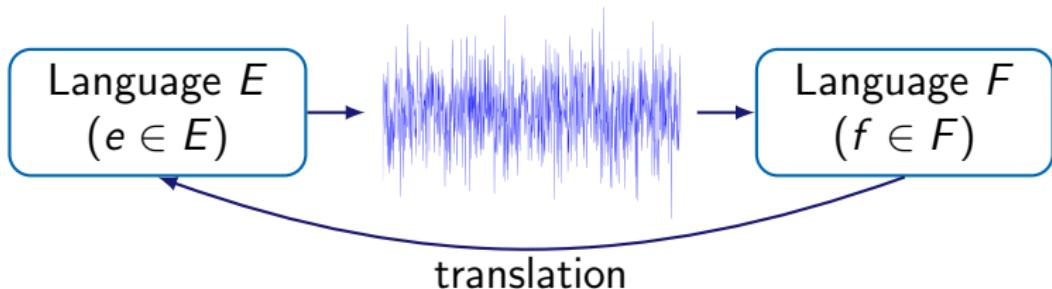


Mathematically:

$$P(e|f)$$

# Statistical Machine Translation

## The Noisy Channel



Mathematically:

$$P(e|f) = \frac{P(e) P(f|e)}{P(f)}$$

$$T(f) = \hat{e} = \operatorname{argmax}_e P(e|f) = \operatorname{argmax}_e P(e) P(f|e)$$

# Statistical Machine Translation

## *Basic equation*

$$T(f) = \hat{e} = \operatorname{argmax}_e P(e) P(f|e)$$

## Language Model

- Takes care of fluency in the target language
- Data: corpora in the target language

# Statistical Machine Translation

## *Basic equation*

$$T(f) = \hat{e} = \operatorname{argmax}_e P(e) \mathbf{P}(\mathbf{f}|e)$$

### Language Model

- Takes care of fluency in the target language
- Data: corpora in the target language

### Translation Model

- Lexical correspondence between languages
- Data: aligned corpora in source and target languages

# Statistical Machine Translation

## Basic equation

$$T(f) = \hat{e} = \operatorname{argmax}_e P(e) P(f|e)$$

### Language Model

- Takes care of fluency in the target language
- Data: corpora in the target language

### Translation Model

- Lexical correspondence between languages
- Data: aligned corpora in source and target languages

### **argmax**

- Search done by the *decoder*

# Statistical Machine Translation

*Remember from the LM lecture...*

$$T(f) = \hat{e} = \operatorname{argmax}_e \mathbf{P}(e) P(f|e)$$

The language model assigns a probability  $P(e)$  to a sequence of words  $e \Rightarrow \{w_1, \dots, w_m\}$ .

$$P(w_1, \dots, w_m) = \prod_{i=1}^m P(w_i | w_{i-(n-1)}, \dots, w_{i-1})$$

- Frequencies in a corpus
- *Softmaxed* output of a neural network

# Statistical Machine Translation

*The translation model  $P(f|e)$*

$$T(f) = \hat{e} = \operatorname{argmax}_e P(e) \mathbf{P}(\mathbf{f}|\mathbf{e})$$

Estimation of the lexical correspondence between languages.

# Statistical Machine Translation

*The translation model  $P(f|e)$*

$$T(f) = \hat{e} = \operatorname{argmax}_e P(e) \mathbf{P}(\mathbf{fle})$$

Estimation of the lexical correspondence between languages.

How can be  $P(f|e)$  characterised?

NULL      Quan      tornes      a      casa      ?

When are you coming back home ?

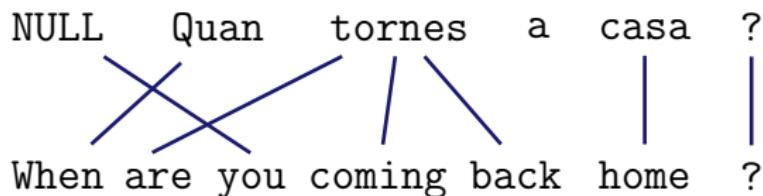
# Statistical Machine Translation

*The translation model  $P(f|e)$*

$$T(f) = \hat{e} = \operatorname{argmax}_e P(e) \mathbf{P}(\mathbf{f}|e)$$

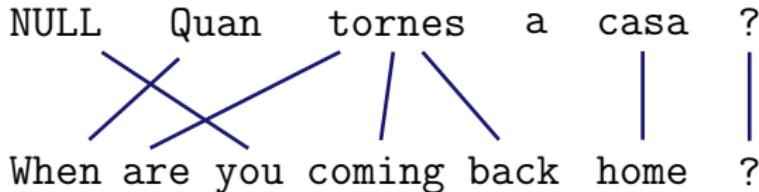
Estimation of the lexical correspondence between languages.

How can be  $P(f|e)$  characterised?



# Statistical Machine Translation

*The translation model  $P(f|e)$*



One should at least model for *each word* in the source language:

- Its translation,
- the number of necessary words in the target language,
- the position of the translation within the sentence,
- and, besides, the number of words that need to be generated from scratch.

# Statistical Machine Translation

## *The translation model $P(f|e)$*

### Word-based models: the IBM models

They characterise  $P(f|e)$  with 4 parameters:  $t$ ,  $n$ ,  $d$  and  $p_1$ .

- Lexical probability  $t$

$t(\text{Quan}|\text{When})$ : the prob. that Quan translates into When.

- Fertility  $n$

$n(3|\text{tornes})$ : the prob. that tornes generates 3 words.

# Statistical Machine Translation

## *The translation model $P(f|e)$*

### Word-based models: the IBM models

They characterise  $P(f|e)$  with 4 parameters:  $t$ ,  $n$ ,  $d$  and  $p_1$ .

- Distortion  $d$

$d(j|i, m, n)$ : the prob. that the word in the  $j$  position generates a word in the  $i$  position.  $m$  and  $n$  are the length of the source and target sentences.

- Probability  $p_1$

$p(\text{you}|\text{NULL})$ : the prob. that the spurious word  $\text{you}$  is generated (from  $\text{NULL}$ ).

# Statistical Machine Translation

## *IBM models*

Back to the example:

NULL Quan tornes a casa ?

# Statistical Machine Translation

## IBM models

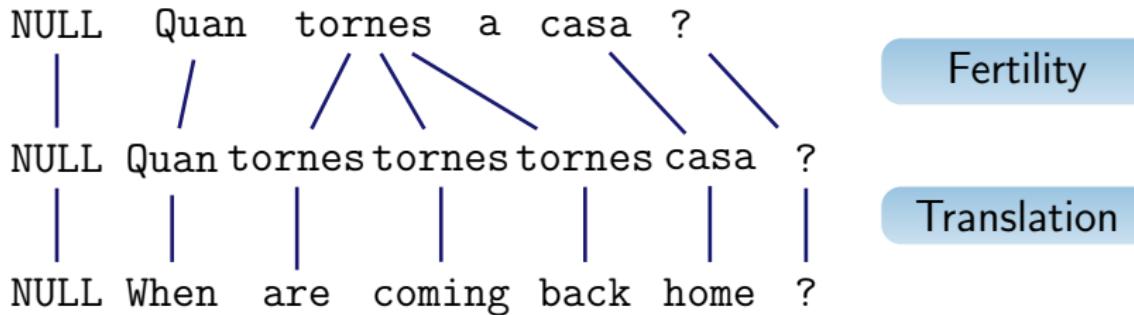
Back to the example:



# Statistical Machine Translation

## IBM models

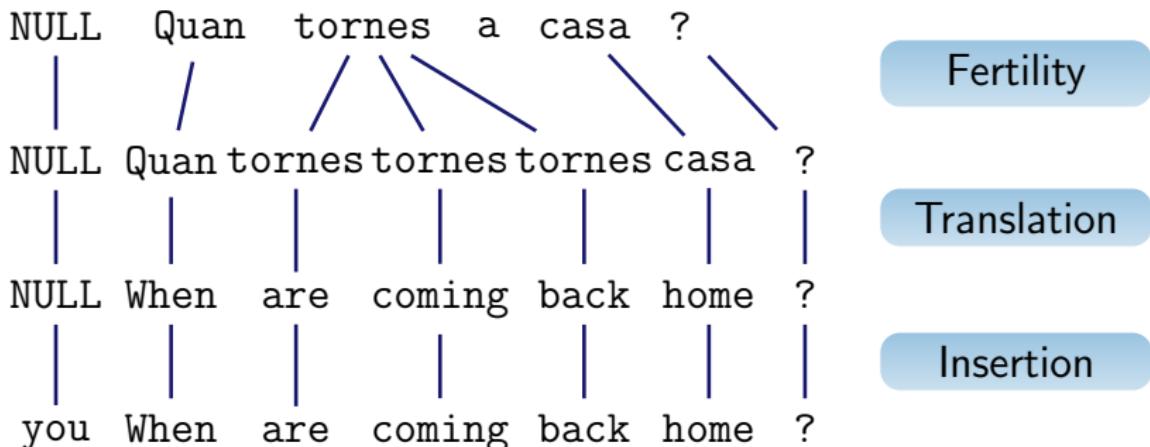
Back to the example:



# Statistical Machine Translation

## IBM models

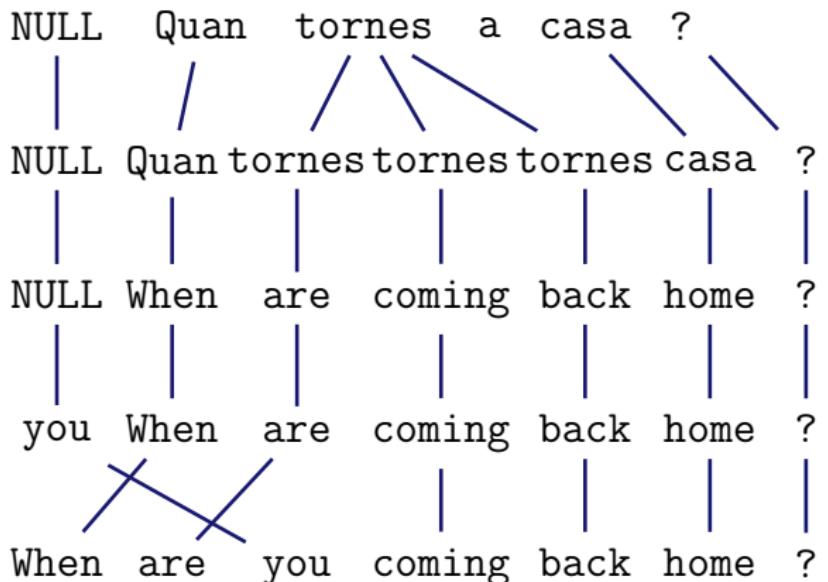
Back to the example:



# Statistical Machine Translation

## IBM models

Back to the example:



Fertility

Translation

Insertion

Distortion

# Statistical Machine Translation

## Decoding

$$T(f) = \hat{e} = \operatorname{argmax}_e P(e) P(f|e)$$

Responsible for the search in the space of possible translations.

Given a model (LM+TM+...), the decoder constructs the possible translations and looks for the most probable one.

# Statistical Machine Translation

## Decoding

$$T(f) = \hat{e} = \operatorname{argmax}_e P(e) P(f|e)$$

Responsible for the search in the space of possible translations.

Given a model (LM+TM+...), the decoder constructs the possible translations and looks for the most probable one.

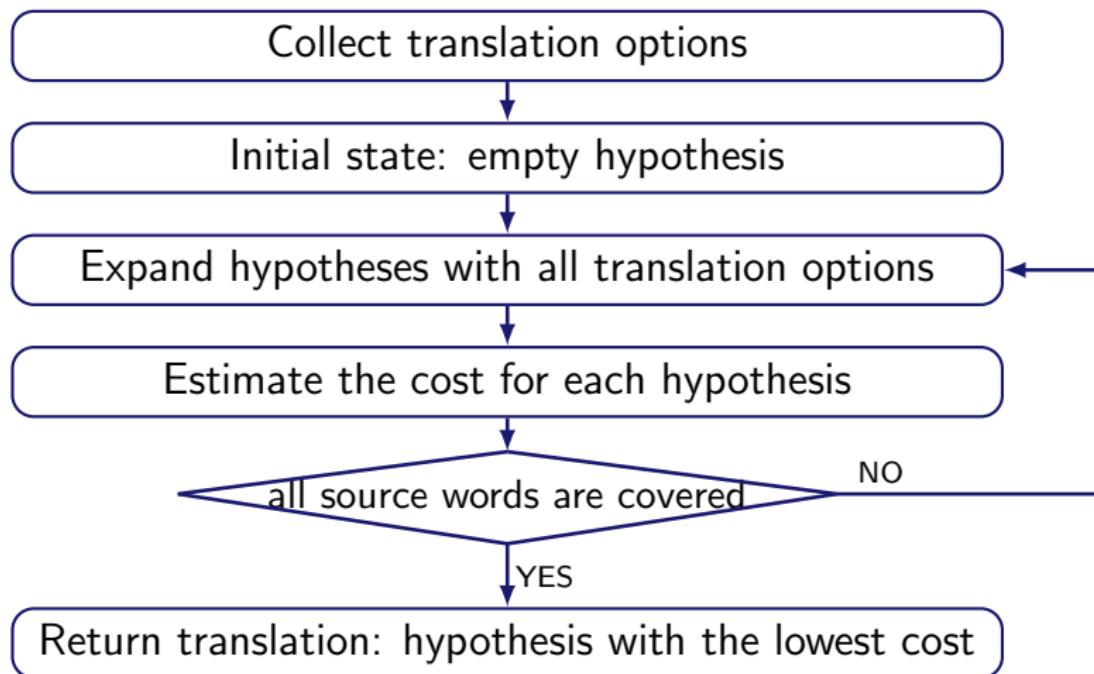
In our context, one can find:

- Greedy decoders. Initial hypothesis (word by word translation) refined iteratively using hill-climbing heuristics.
- Beam search decoders.

# Statistical Machine Translation

## Beam Search Decoder

### Core algorithm



# Statistical Machine Translation

*State-of-the-art*

- Log-linear models
- Phrase-based extension

$$\hat{e} = \operatorname{argmax}_e \log P(e|f) = \operatorname{argmax}_e \sum \lambda_m h_m(f, e)$$

# Neural Machine Translation

1 Introduction

2 Statistical Machine Translation

3 Neural Machine Translation

- RNN Architecture
- Transformer Architecture

4 Software & References

### The Encoder–Decoder Model (with attention)

- 1 encodes a sequence of word vectors into a fixed-sized context vector
- 2 decodes the fixed-sized vector back into a variable-length sequence

### The Encoder–Decoder Model (with attention)

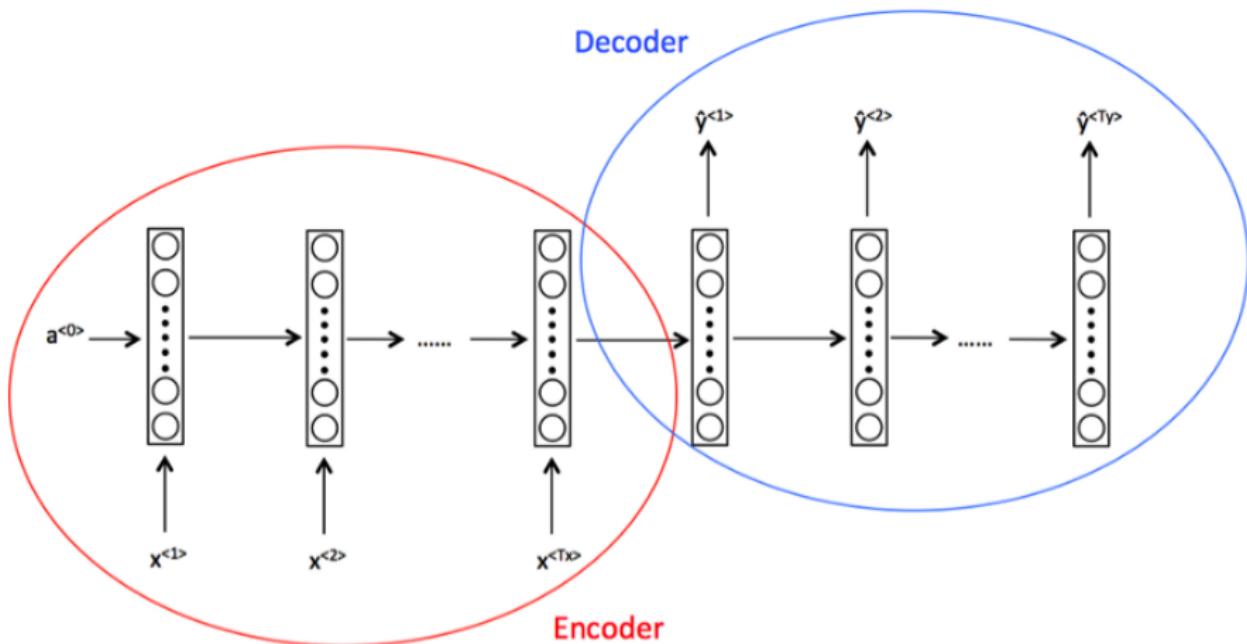
- 1 encodes a sequence of word vectors into a fixed-sized context vector
- 2 decodes the fixed-sized vector back into a variable-length sequence

Several NLP tasks use nowadays enc–dec architectures:

- Machine translation, but also...
- text summarisation, question answering, chatbots, speech recognition...

# Neural Machine Translation

*Remember from the Sequence Modelling lecture...*



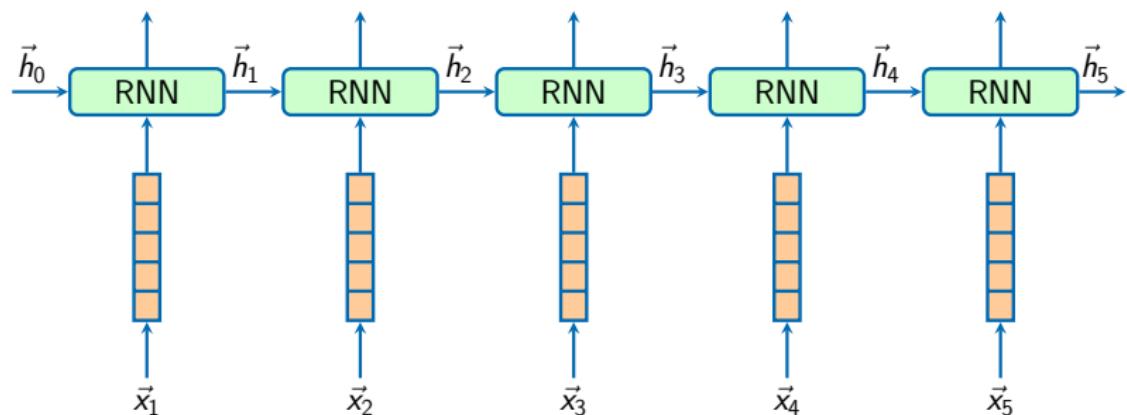
<https://medium.com/machine-learning-bites/>

/deeplearning-series-sequence-models-7855babeb586

# Vanilla RNN Architecture

## Encoder

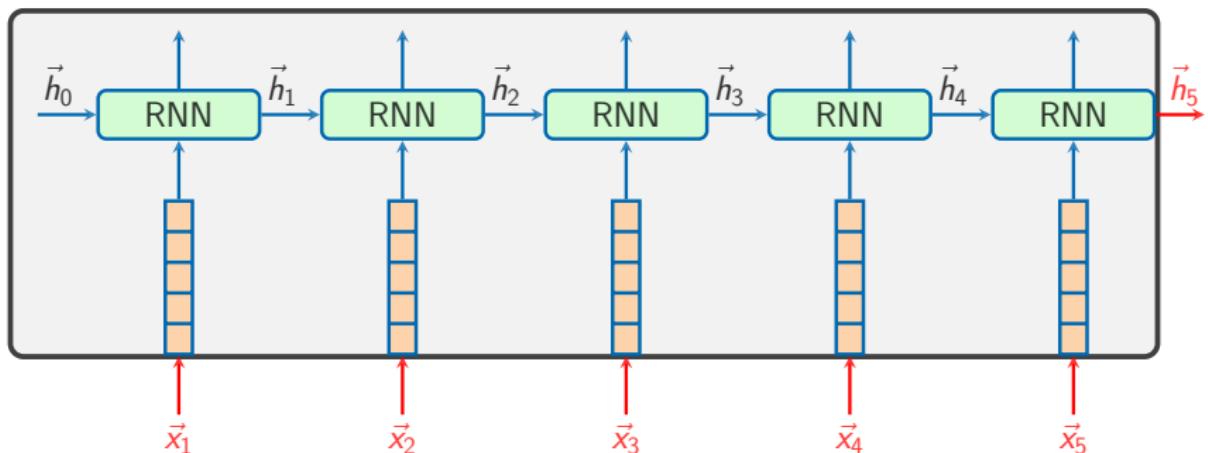
- Sequence encoded with RNNs
- Each word is a *time step*



# Vanilla RNN Architecture

## Encoder

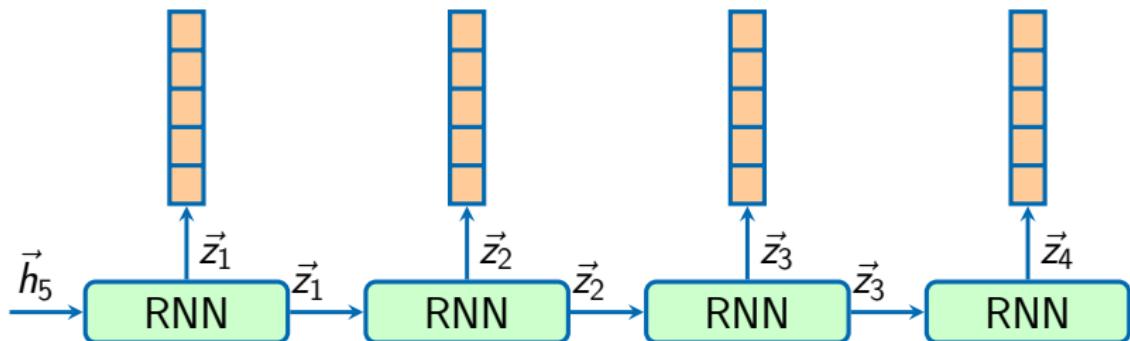
- Sequence encoded with RNNs
- Each word is a *time step*



# Vanilla RNN Architecture

## Decoder

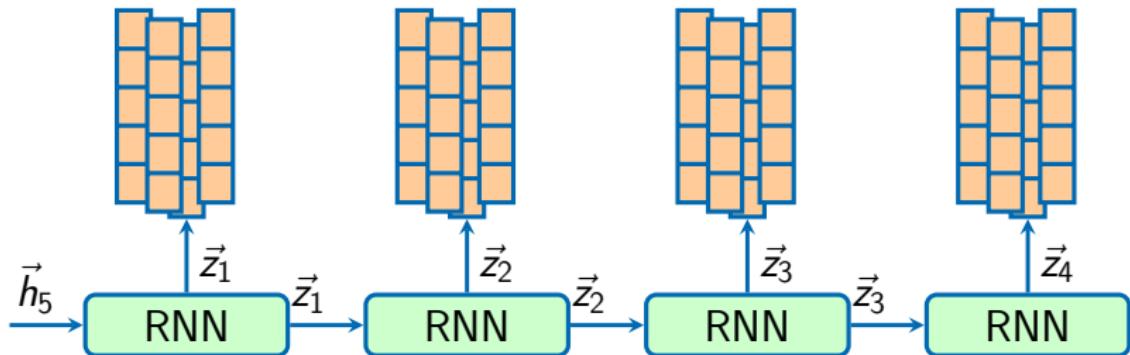
- Inverted encoder with special characteristics



# Vanilla RNN Architecture

## Decoder

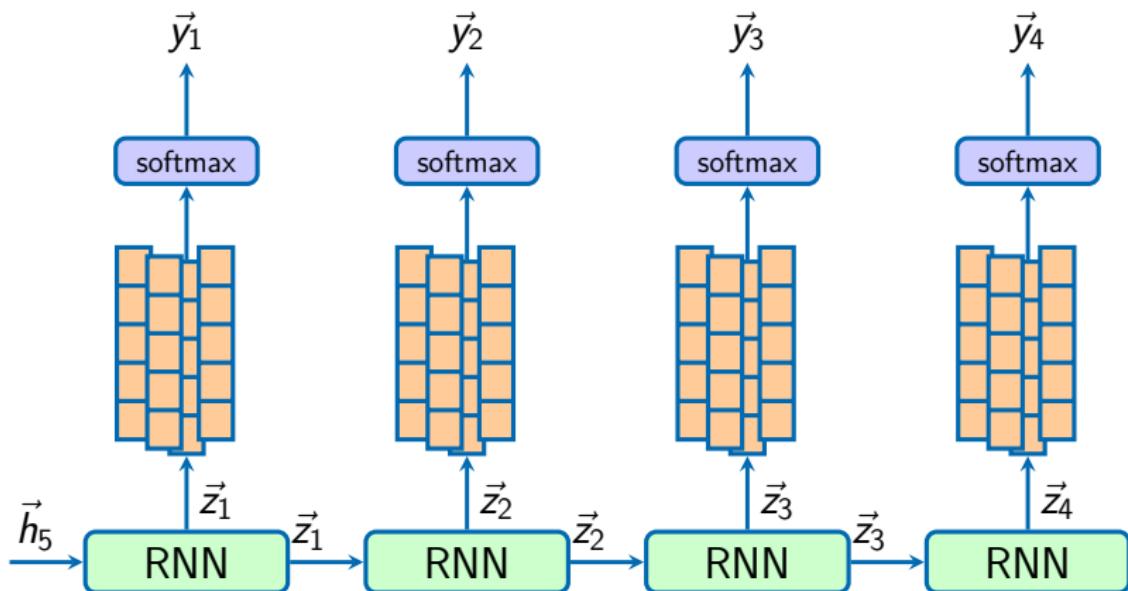
- Inverted encoder with special characteristics



# Vanilla RNN Architecture

## Decoder

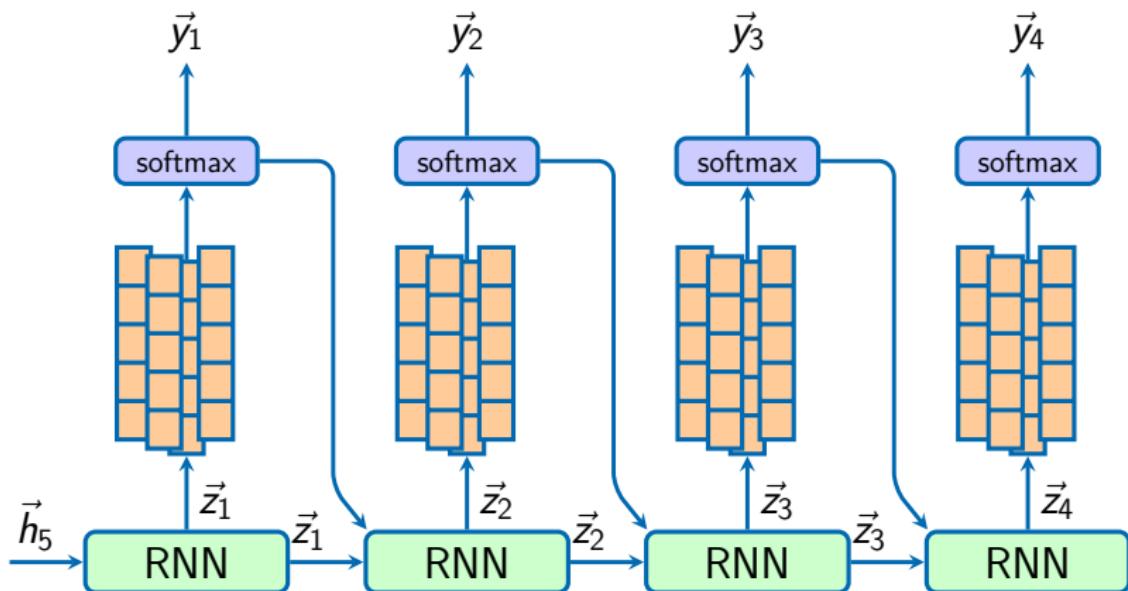
- Inverted encoder with special characteristics



# Vanilla RNN Architecture

## Decoder

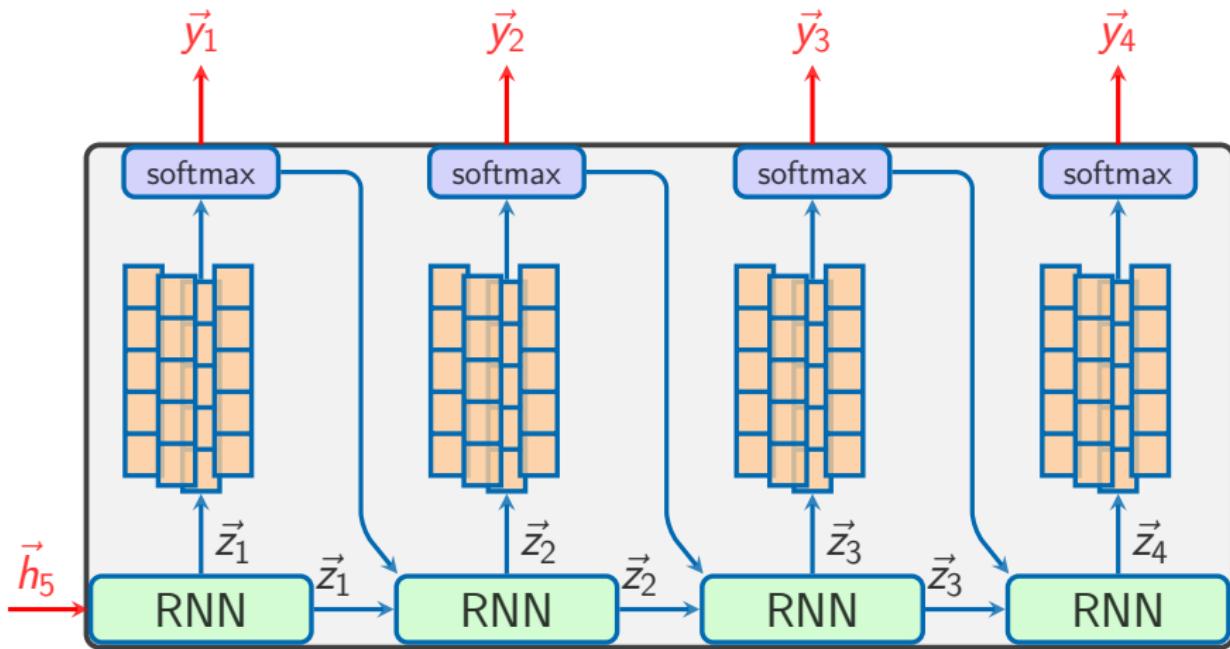
- Inverted encoder with special characteristics



# Vanilla RNN Architecture

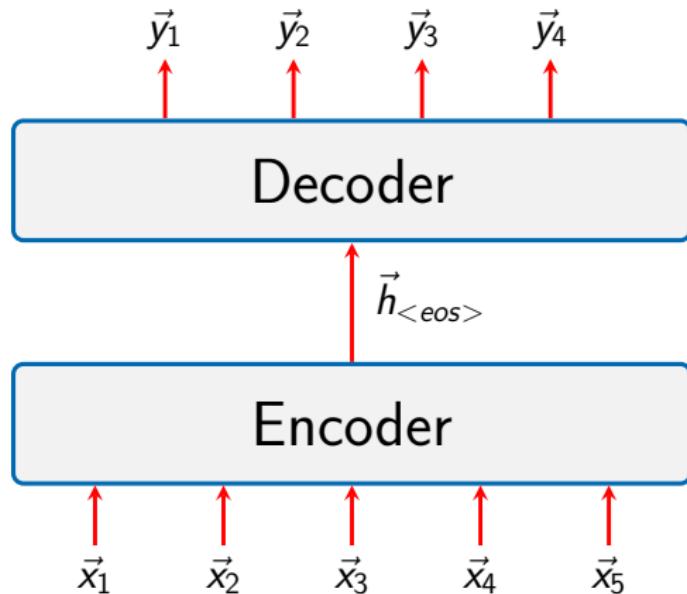
## Decoder

- Inverted encoder with special characteristics



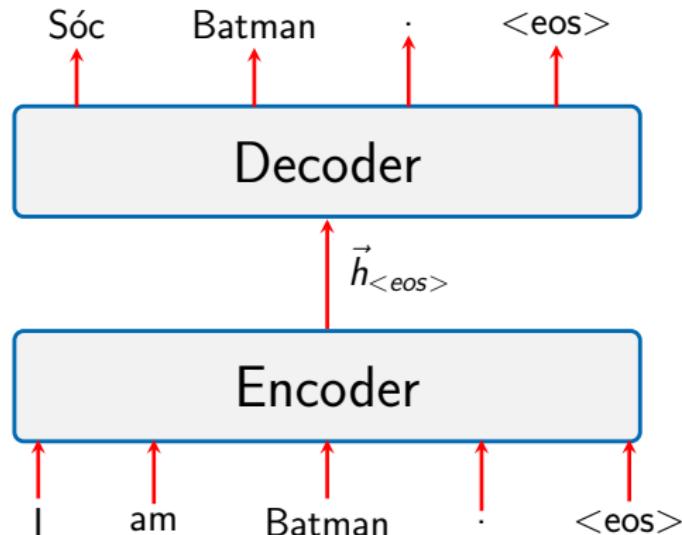
# Vanilla RNN Architecture

## Encoder–Decoder



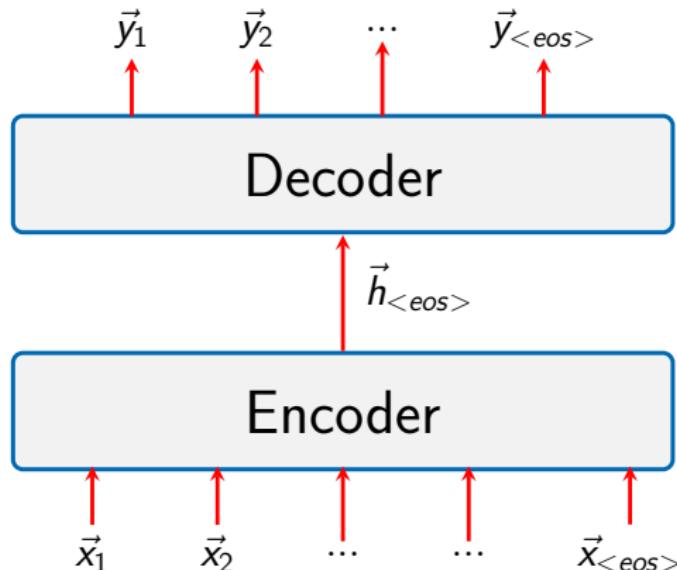
# Vanilla RNN Architecture

## Encoder–Decoder



# Vanilla RNN Architecture

## Encoder–Decoder

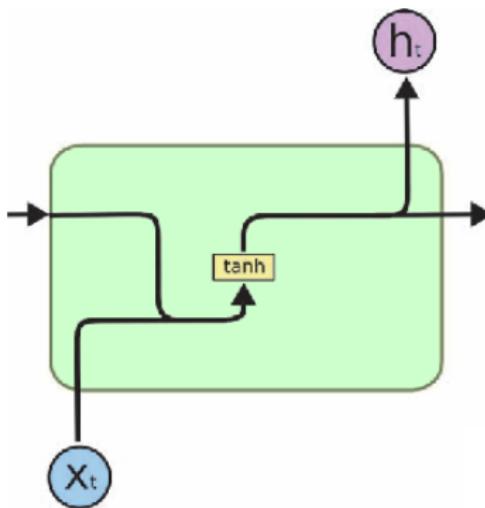


# Vanilla RNN Architecture

## RNN cells

- RNN cells can be any kind

**tanh**

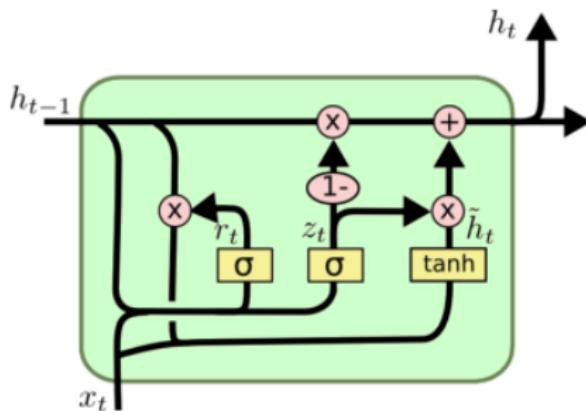


# Vanilla RNN Architecture

## RNN cells

- RNN cells can be any kind

### GRU, Gated Recurrent Unit

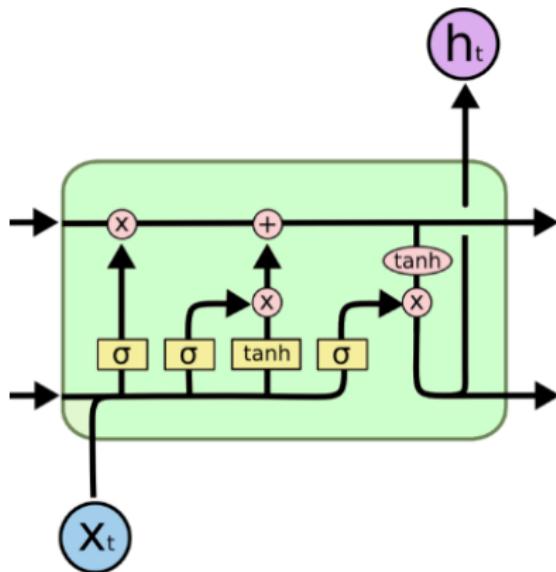


# Vanilla RNN Architecture

## RNN cells

- RNN cells can be any kind

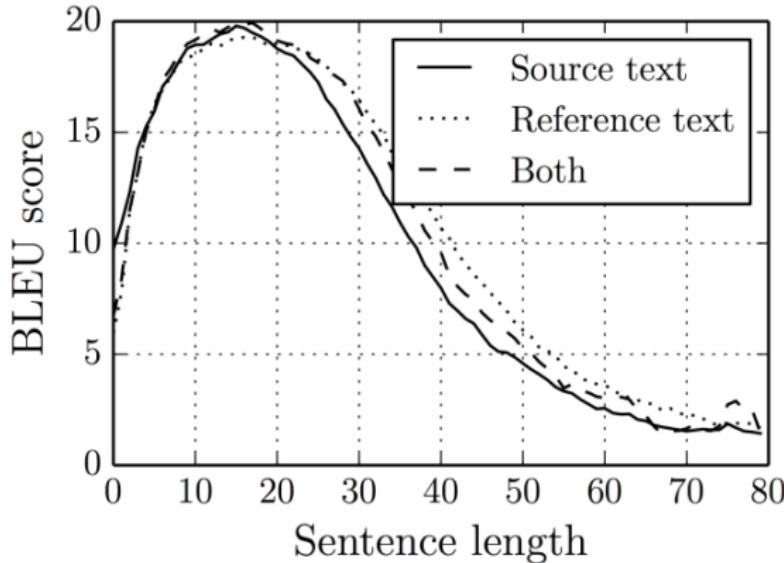
### LSTM, Long Short Term Memory Networks



# Vanilla RNN Architecture

## Why Vanilla?

- Performance drops with long sentences

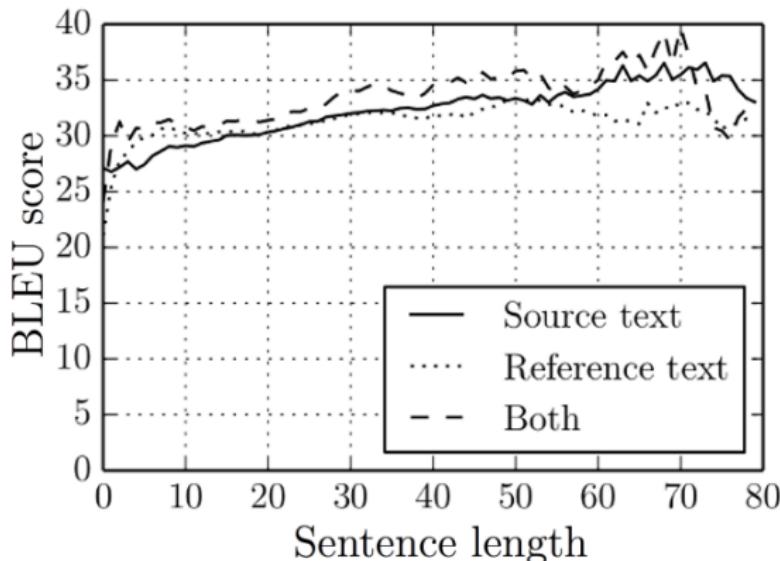


[Cho et al., 2014]

# Vanilla RNN Architecture

## Why Vanilla?

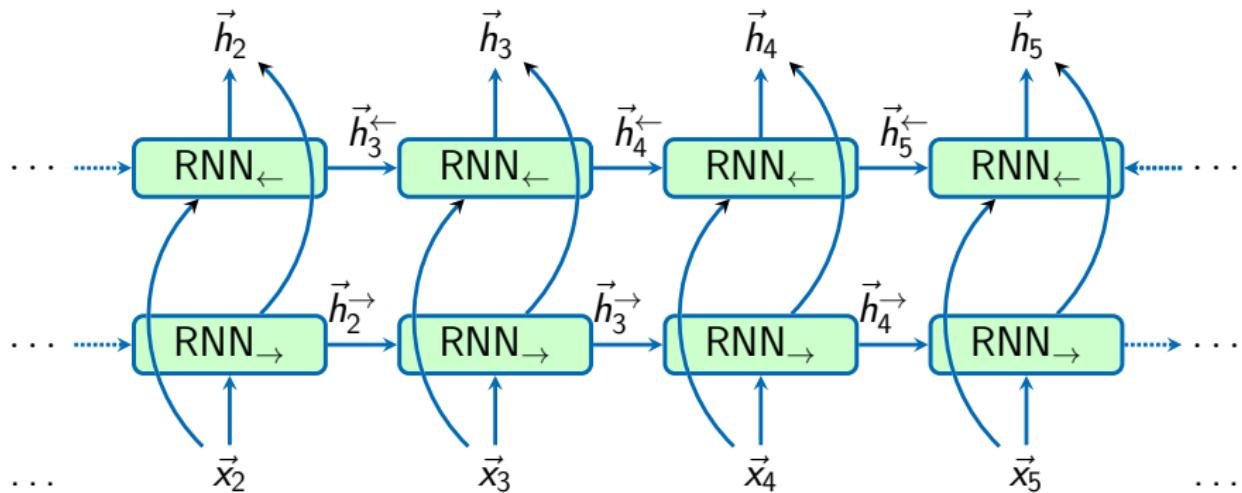
- SMT was better at that!



[Cho et al., 2014]

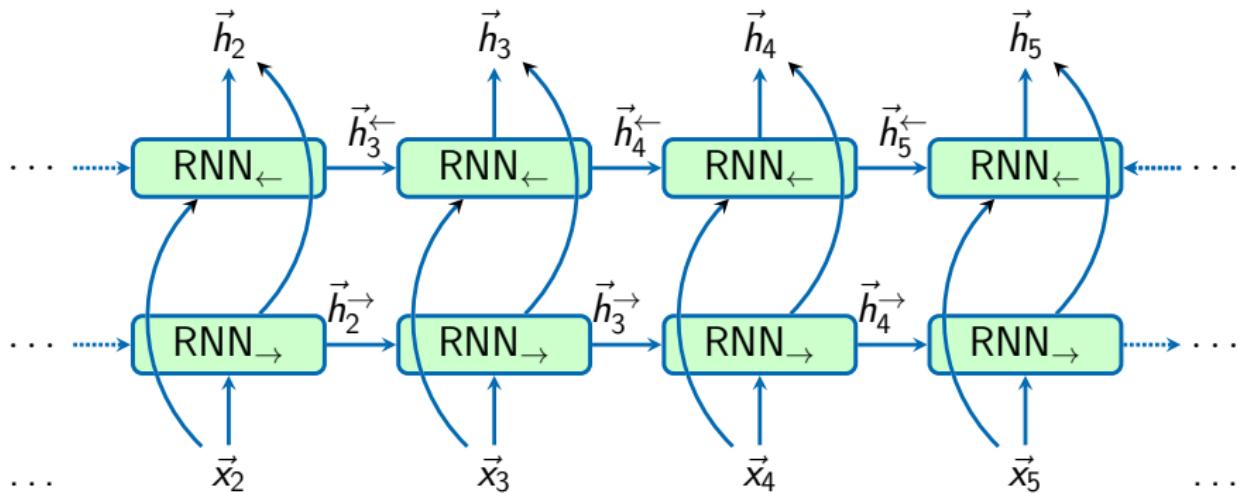
# RNN Architecture

## Bidirectional Encoder Hidden Layer



# RNN Architecture

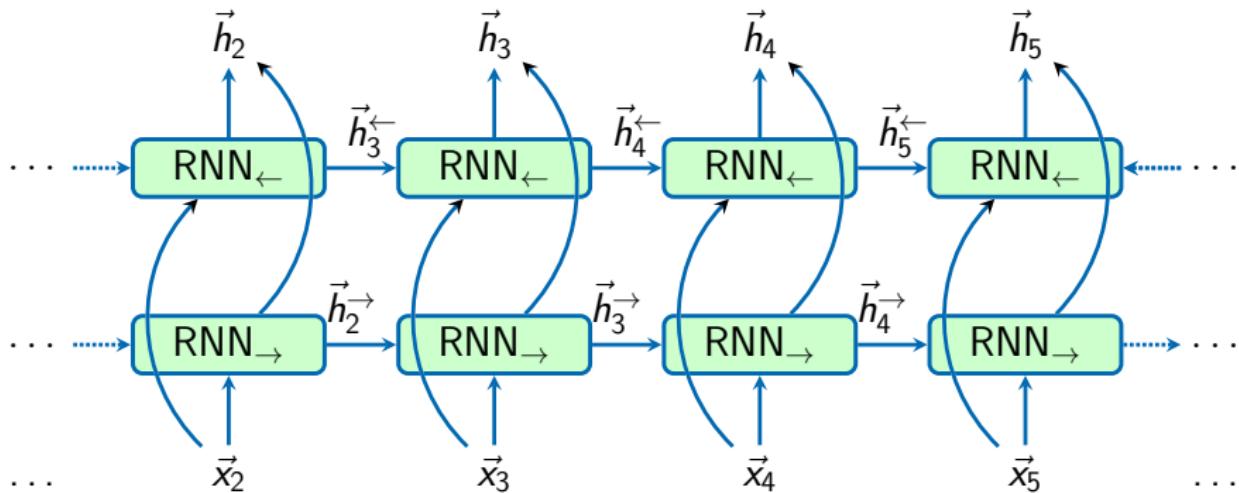
## Bidirectional Encoder Hidden Layer



- Helps but not enough
- What else?

# RNN Architecture

## Bidirectional Encoder Hidden Layer

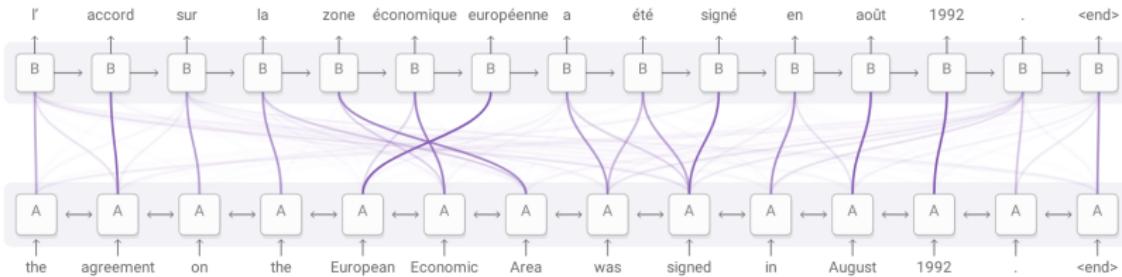


- Helps but not enough
- What else? **Attention!**

# RNN Architecture

## The Attention Mechanism

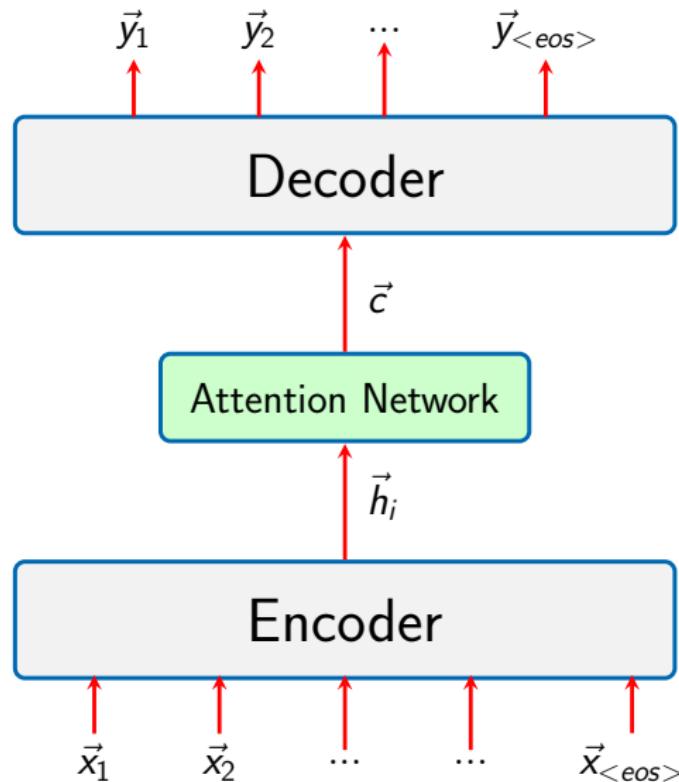
- Intuition: Not all the words contribute equally for a translation
- Let's weight! (weights, softmaxs, nns...)



<https://distill.pub/2016/augmented-rnns>

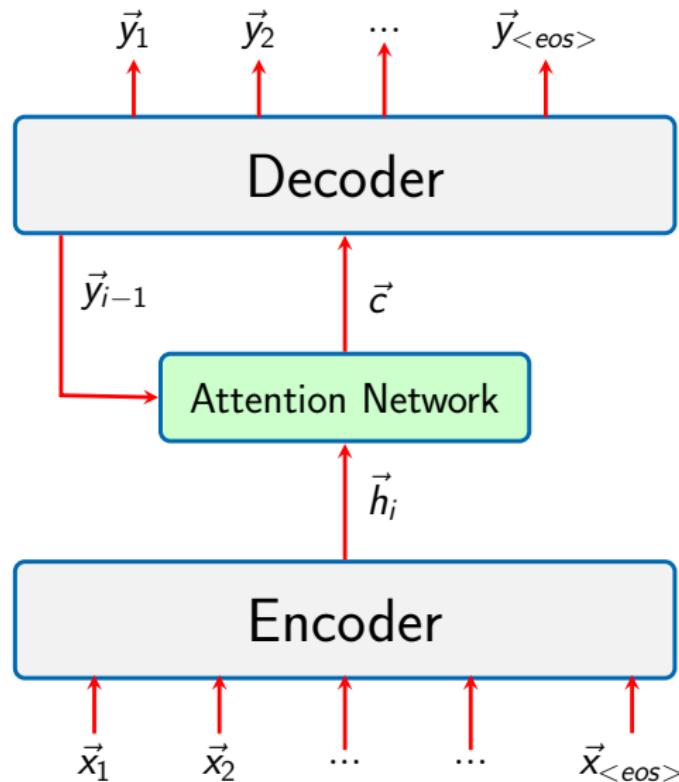
# RNN Architecture

## The Attention Mechanism



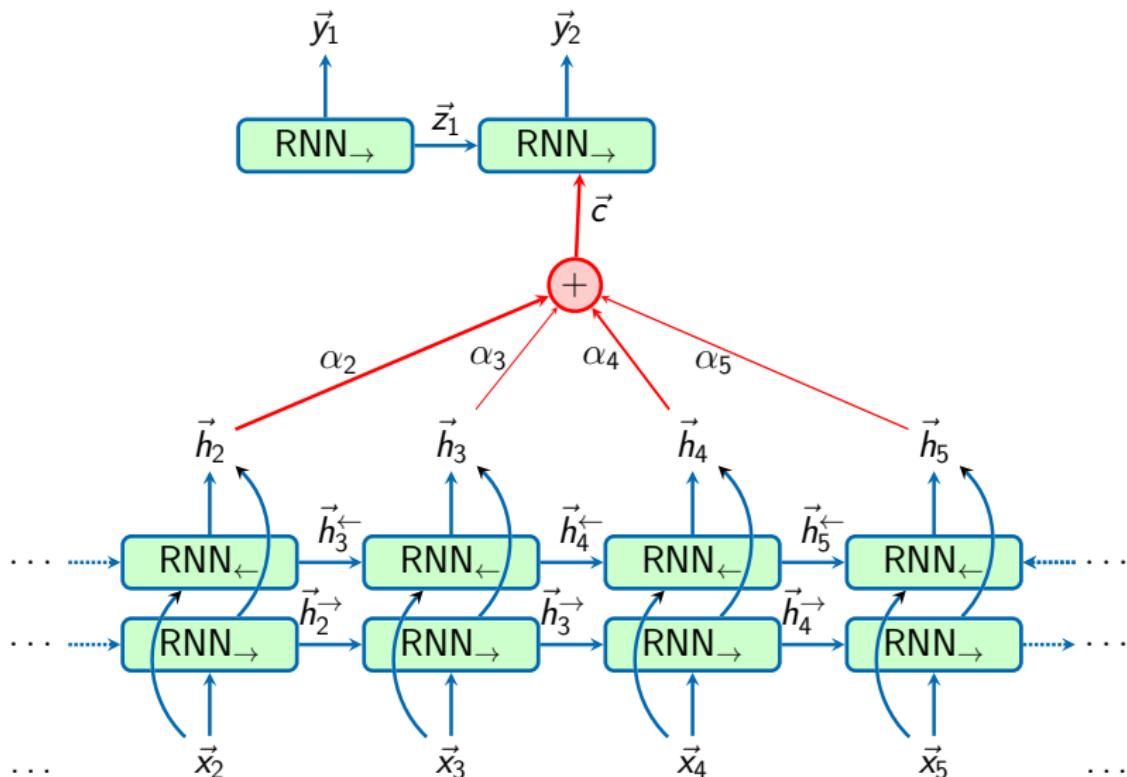
# RNN Architecture

## The Attention Mechanism



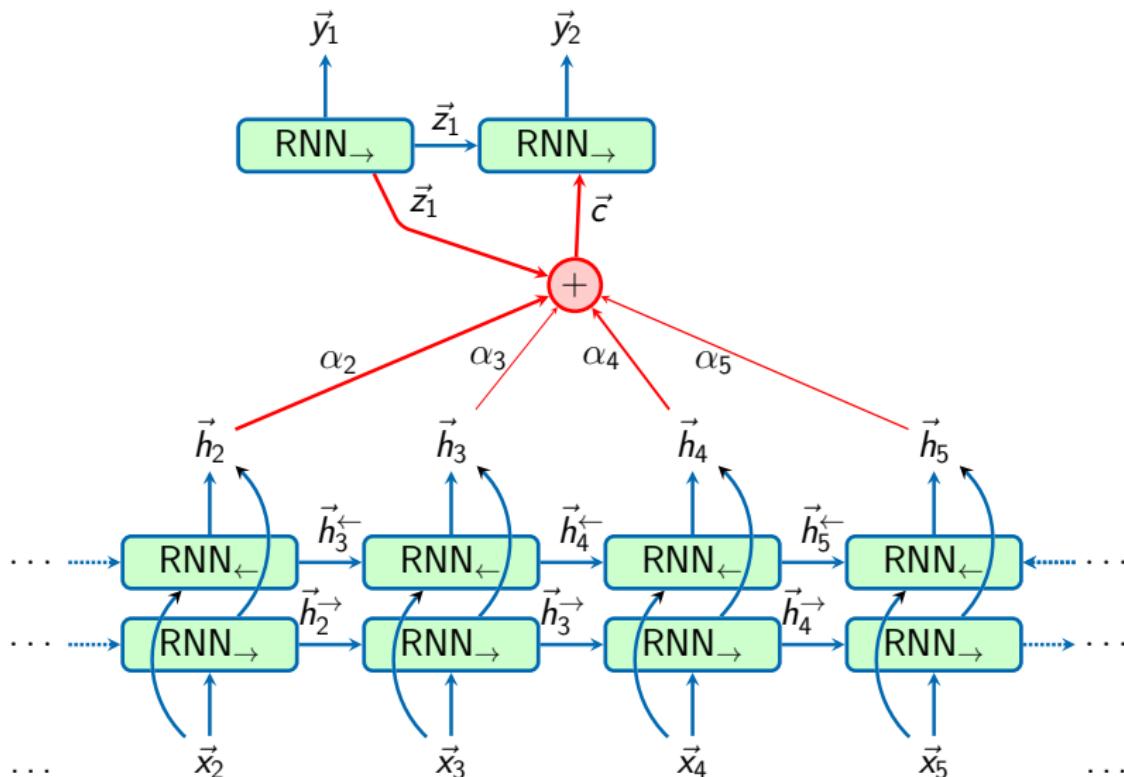
# RNN Architecture

## The Attention Mechanism



# RNN Architecture

## The Attention Mechanism



# RNN Architecture

## The Attention Mechanism

### Encoder with bidirectional RNNs

Hidden vectors: concatenation of the forward ( $\overrightarrow{\mathbf{h}}_i$ ) and backward ( $\overleftarrow{\mathbf{h}}_i$ ) encoder RNN hidden states:

$$\begin{aligned}\mathbf{h}_i &= \left[ \overleftarrow{\mathbf{h}}_i, \overrightarrow{\mathbf{h}}_i \right] \\ &= \left[ f(\overleftarrow{\mathbf{h}}_{i-1}, \mathbf{r}_i), f(\overrightarrow{\mathbf{h}}_{i+1}, \mathbf{r}_i) \right],\end{aligned}$$

- $f$  is a recurrent unit (LSTM, GRU, etc.)
- $\mathbf{r}_i$  is the embedding space representation of the source word at position  $i$ :  $\mathbf{r}_i = \mathbf{W}_x \cdot \mathbf{x}_i$ .

# RNN Architecture

## The Attention Mechanism

### Decoder with RNNs

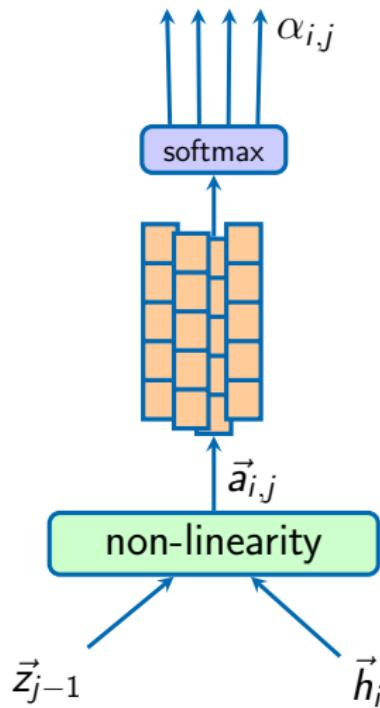
$$\mathbf{z}_j = g(\mathbf{z}_{j-1}, \mathbf{t}_{j-1}, \mathbf{q}_j)$$

$$\mathbf{t}_{j-1} = \mathbf{W}_y \cdot \mathbf{y}_{j-1},$$

- $\mathbf{z}_j$  is the recurrent hidden state of the decoder
- $\mathbf{t}_j$  is the continuous representation of the target word
- $g$  is a non-linear function
- $\mathbf{W}_y$  is the matrix of the target embeddings
- The weighted context vector  $\mathbf{q}_j$  is calculated by the *attention mechanism*

# RNN Architecture

## Attention Network



# RNN Architecture

## Attention Network

In equations:

$$a(\mathbf{z}_{j-1}, \mathbf{h}_i) = \mathbf{v}_a \cdot \tanh(\mathbf{W}_a \cdot \mathbf{z}_{j-1} + \mathbf{U}_a \cdot \mathbf{h}_i)$$

$$\alpha_{ij} = \text{softmax}(a(\mathbf{z}_{j-1}, \mathbf{h}_i))$$

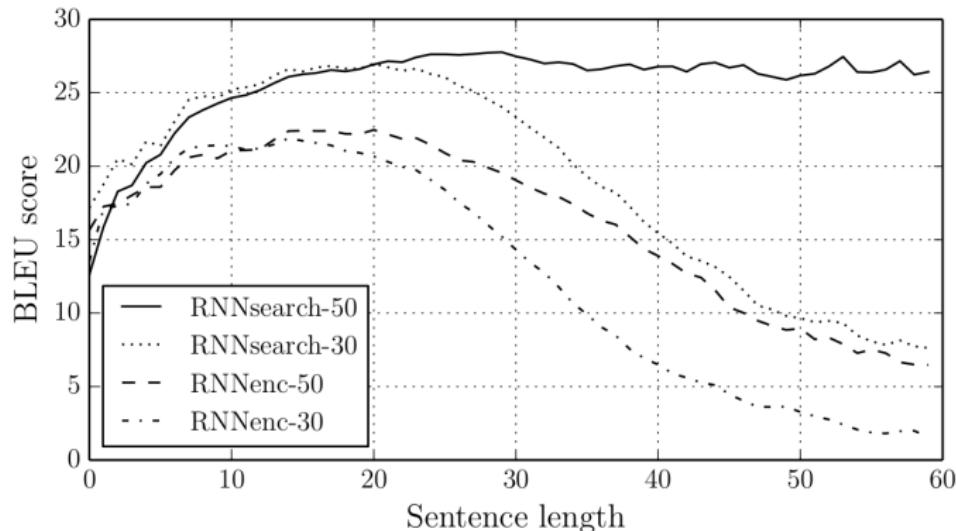
$$\mathbf{q}_j = \sum_i \alpha_{ij} \mathbf{h}_i$$

The attention mechanism takes the decoder's previous hidden state  $\mathbf{z}_{j-1}$  and the context vector  $\mathbf{h}_i$  as inputs and weighs them up with the trainable weight matrices  $\mathbf{W}_a$  and  $\mathbf{U}_a$ , respectively.

# RNN Architecture

## *The Attention Mechanism*

- What achieves attention?



[Bahdanau et al., 2015]

# Transformer Architecture

*So, Attention is all you Need?!*

## Attention Is All You Need

Ashish Vaswani\*  
Google Brain  
avaswani@google.com

Noam Shazeer\*  
Google Brain  
noam@google.com

Niki Parmar\*  
Google Research  
nikip@google.com

Jakob Uszkoreit\*  
Google Research  
usz@google.com

Llion Jones\*  
Google Research  
llion@google.com

Aidan N. Gomez\* †  
University of Toronto  
aidan@cs.toronto.edu

Lukasz Kaiser\*  
Google Brain  
lukaszkaiser@google.com

Illia Polosukhin\* ‡  
illia.polosukhin@gmail.com

### Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions

# Transformer Architecture

*Remember the Language Modelling Lecture*

We have already done the hard work:

- Remember self-attention?
- Remember multi-head attention?

<http://jalammar.github.io/illustrated-transformer/>

# Transformer Architecture

## *Self-Attention*

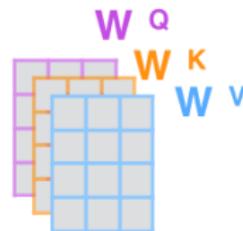
1) This is our input sentence

2) We embed each word

3) We multiply  $X$  with weight matrices

4) Calculate attention using the resulting  $Q/K/V$  matrices

Thinking  
Machines



$$\text{softmax} \left( \frac{\begin{pmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{pmatrix} \times \begin{pmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{pmatrix}}{\sqrt{d_k}} \right) \begin{pmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{pmatrix} = \begin{pmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{pmatrix}$$

Diagram illustrating the calculation of the final output  $Z$  from the query  $Q$  and key  $K^T$  matrices. The diagram shows the softmax function applied to the product of  $Q$  and  $K^T$  scaled by  $\sqrt{d_k}$ , followed by multiplication with the value matrix  $V$  to produce the output  $Z$ .

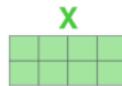
# Transformer Architecture

## Multi-head Attention

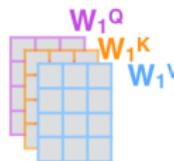
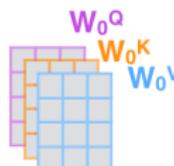
1) This is our input sentence

Thinking Machines

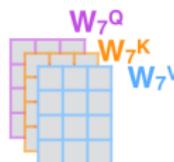
2) We embed each word



3) Split into 8 heads.  
We multiply  $X$  with weight matrices



...



4) Calculate attention using the resulting  $Q/K/V$  matrices



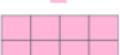
5) Concatenate the resulting  $Z$  matrices, then multiply with weight matrix  $W^o$  to produce the output of the layer



$W^o$



$Z$



# Transformer Architecture

## *Structure*

- (Self)Attention has multiple heads
- Encoder has multiple encoder layers
- Decoder has multiple decoder layers
- (Inter)Attention behaves as usual (one per decoder layer)

# Transformer Architecture

## *Structure*

- (Self)Attention has multiple heads
- Encoder has multiple encoder layers
- Decoder has multiple decoder layers
- (Inter)Attention behaves as usual (one per decoder layer)

**Why all this?**

# Transformer Architecture

## *Structure*

- (Self)Attention has multiple heads
- Encoder has multiple encoder layers
- Decoder has multiple decoder layers
- (Inter)Attention behaves as usual (one per decoder layer)

**Why all this?** Because it works...

# Transformer Architecture

## *Structure*

Why all this? Because it works... **but also:**

# Transformer Architecture

## Structure

Why all this? Because it works... **but also:**

- Sequentiality in RNNs prohibits parallelization within instances
- CNNs are easy to parallelise but need  $\log(n)$  layers to consider words at positions  $i$  and  $i + n$

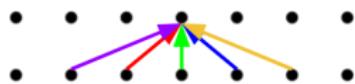
# Transformer Architecture

## Structure

Why all this? Because it works... **but also:**

- Sequentiality in RNNs prohibits parallelization within instances
- CNNs are easy to parallelise but need  $\log(n)$  layers to consider words at positions  $i$  and  $i + n$

Convolution



Self-Attention



# Transformer Architecture

## Structure

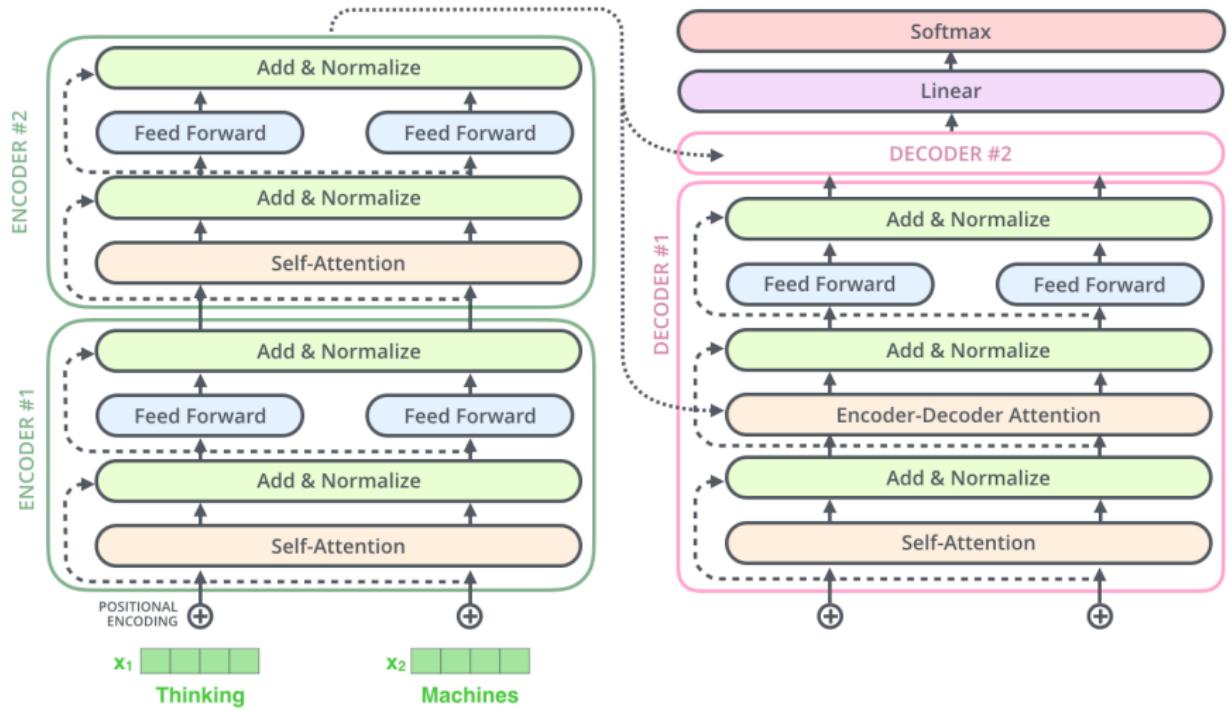


- Convolution: a different linear transformation for each relative position. Allows you to distinguish what information came from where.
- Multiple attention layers (head) in parallel. Each head uses different linear transformations which can learn different relationships.

(Lukasz Kaiser slides)

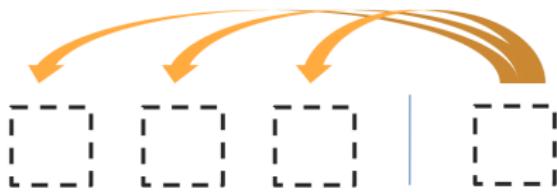
# Transformer Architecture

## The Complete (Reduced) Picture



# Transformer Architecture

## *Different Attentions*



Encoder-Decoder Attention



Encoder Self-Attention

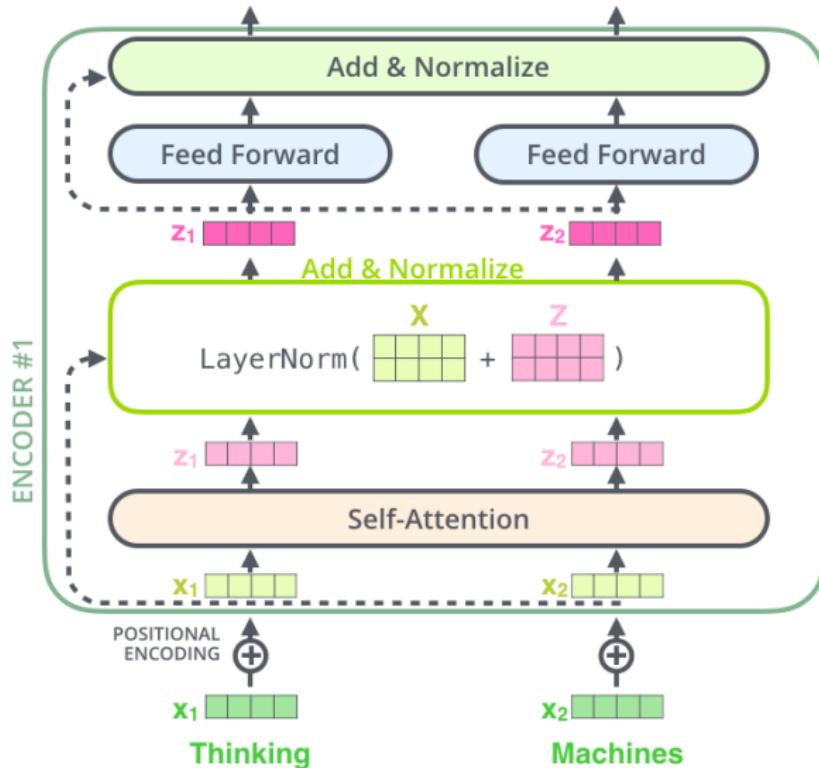


Masked Decoder Self-Attention

(Lukasz Kaiser slides)

# Transformer Architecture

## An Encoder Layer



# Transformer Architecture

## *Positional Embeddings*

- Transformer is not sequential
- No information about the sentence structure, kind of BoWs
- What do we do? Encode position too!
- How? As sinusoidal waves (??)

$$\text{PE}(pos, 2i) = \sin(pos/10000^{2i/d})$$

$$\text{PE}(pos, 2i + 1) = \cos(pos/10000^{2i/d})$$

where  $pos$  is the position of the token,  $i$  is the dimension.

# Transformer Architecture

## Positional Embeddings



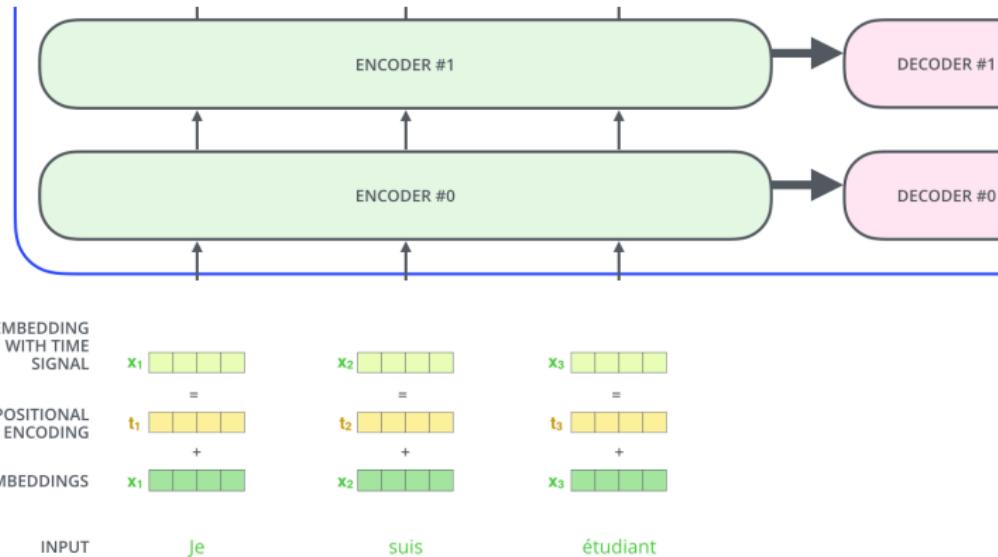
$$\text{PE}(pos, 2i) = \sin(pos/10000^{2i/d})$$

$$\text{PE}(pos, 2i + 1) = \cos(pos/10000^{2i/d})$$

where  $pos$  is the position of the token,  $i$  is the dimension.

# Transformer Architecture

## Positional Embeddings



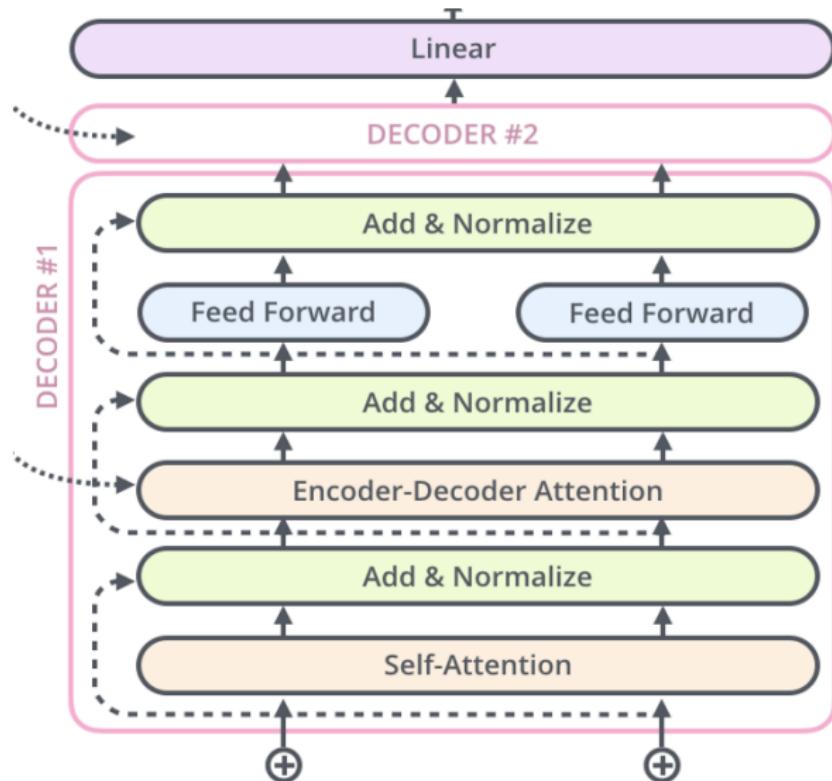
# Transformer Architecture

## Encoder vs. Decoder Layers

- **Masking.** The decoder cannot see *the future* when predicting the next word.
- How?  $-\infty$  before self-attention softmax
- **Enc-Dec Attention.** Queries are taken from the layer below it, but keys and values from the final output of the encoder.
- The decoder adds an additional **linear and softmax** layer (just as RNNs NMT)

# Transformer Architecture

## A Decoder Layer



# Transformer Architecture

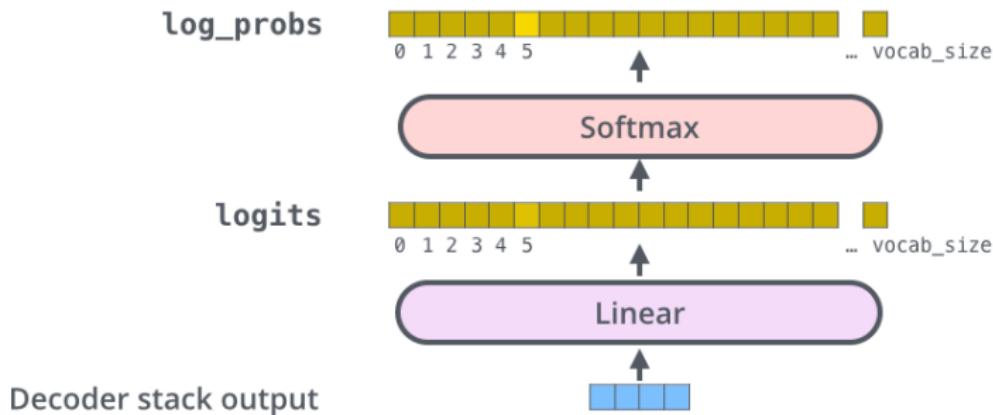
## *The Icing on the Cake*

Which word in our vocabulary  
is associated with this index?

am

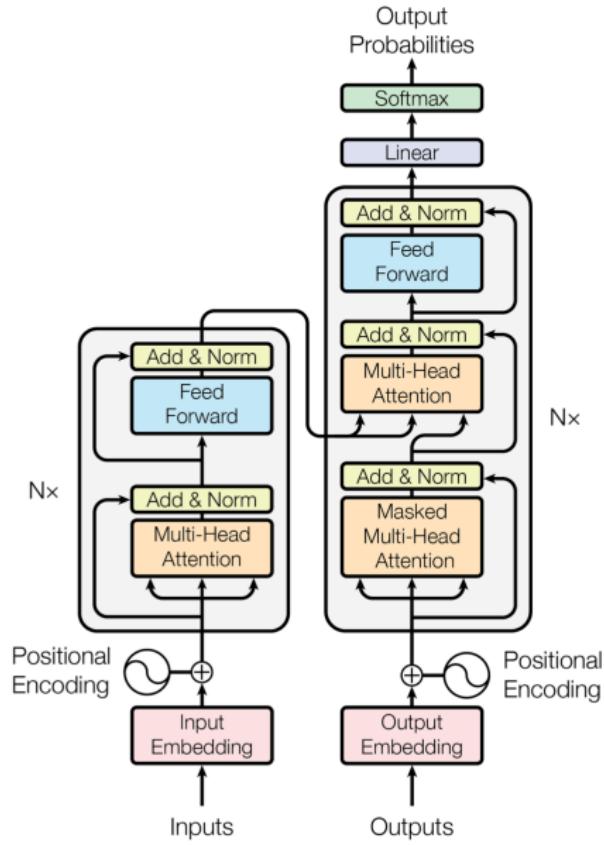
Get the index of the cell  
with the highest value  
(`argmax`)

5



# Transformer Architecture

*The Complete (Original) Picture*



# Transformer Architecture

*Today's Icing on the Cake*

Let's see some animations

<http://jalammar.github.io/illustrated-transformer/>

I love this blog!

# Outline

- 1 Introduction
- 2 Statistical Machine Translation
- 3 Neural Machine Translation
- 4 Software & References

# Software & References

## *NMT Packages*

### ■ Nematus

<https://github.com/EdinburghNLP/nematus>

### ■ Marian

<https://github.com/marian-nmt/marian>

### ■ OpenNMT

<http://opennmt.net/>

### ■ Facebook

<https://github.com/pytorch/fairseq>

# Software & References

## *NMT & SMT Packages*

- **Google**

`https://github.com/tensorflow/tensor2tensor/blob/master/tensor2tensor/models/transformer.py`

- ...

- **Moses (SMT!)**

`https://github.com/moses-smt/mosesdecoder`

# Software & References

## References Used

- Peter F. Brown, Vincent J. Della Pietra, Stephen A. Della Pietra and Robert L. Mercer. 1993. **The mathematics of statistical machine translation: parameter estimation.** Computational Linguistics, 19(2). Pages 263–311.
- Philipp Koehn, Franz Josef Och and Daniel Marcu. **Statistical phrase-based translation.** 2015. Proceedings of the Human Language Technology and North American Association for Computational Linguistics Conference (HLT/NAACL).
- D. Bahdanau, K. Cho and Y. Bengio. 2014. **Neural machine translation by jointly learning to align and translate.** Proceedings of the 3rd Int. Conf. Learn. Represent. Pages 1–15.

# Software & References

## *References Used*

- I. Sutskever, O. Vinyals and Q. V. Le. 2014. **Sequence to sequence learning with neural networks.** Proc. Adv. Neural Inf. Process. Syst. 27. Pages 3104–3112.
- Minh-Thang Luong, Hieu Pham, and Christopher D. Manning. 2015. **Effective Approaches to Attention-based Neural Machine Translation.** Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing. Pages 1412–1421.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser and Illia Polosukhin. 2017. **Attention Is All You Need.** 31st Conference on Neural Information Processing Systems (NIPS 2017).

Too much for a day?

Questions?

# Machine Translation

Cristina España-Bonet

UdS & DFKI, Saarbrücken, Germany

Artificial Intelligence with Deep Learning

2nd May 2019