

# AIDL: PROJECTS

• • •

Session 1 (2019/03/05): Introduction

# Issey Masuda Mora

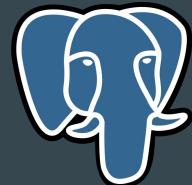
Education: Telecommunications Engineer



Currently working: Deep learning team lead at Vilynx



Technologies:



# Hand-raised poll

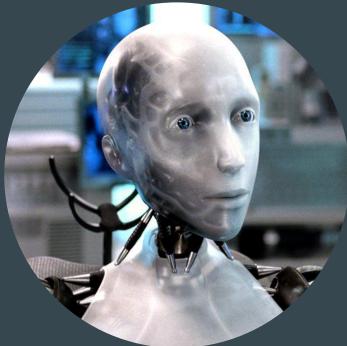
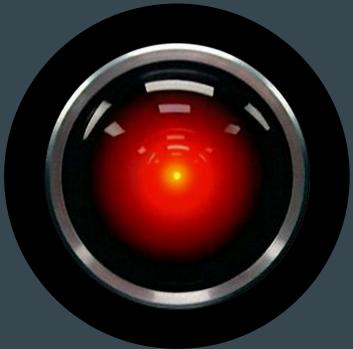
- Programming skills?
- Working on software development? Or management?

# Subject goals

**“Give them the tools to  
apply DL in the industry”**



# Expectations: what not to expect



# Expectations

- Implement NN models
- Ready to prod
- Useful for current job
- Product MVP
- Solve real-world problems
- State of the art

# Subject goals

**Product development**

**Ready to prod model**

**Implement DL model**

# Subject goals

Tangibles:

- Implement DL models
- Optimize resource usage / increase efficiency
- Optimize models
- Deploy to production

Intangibles:

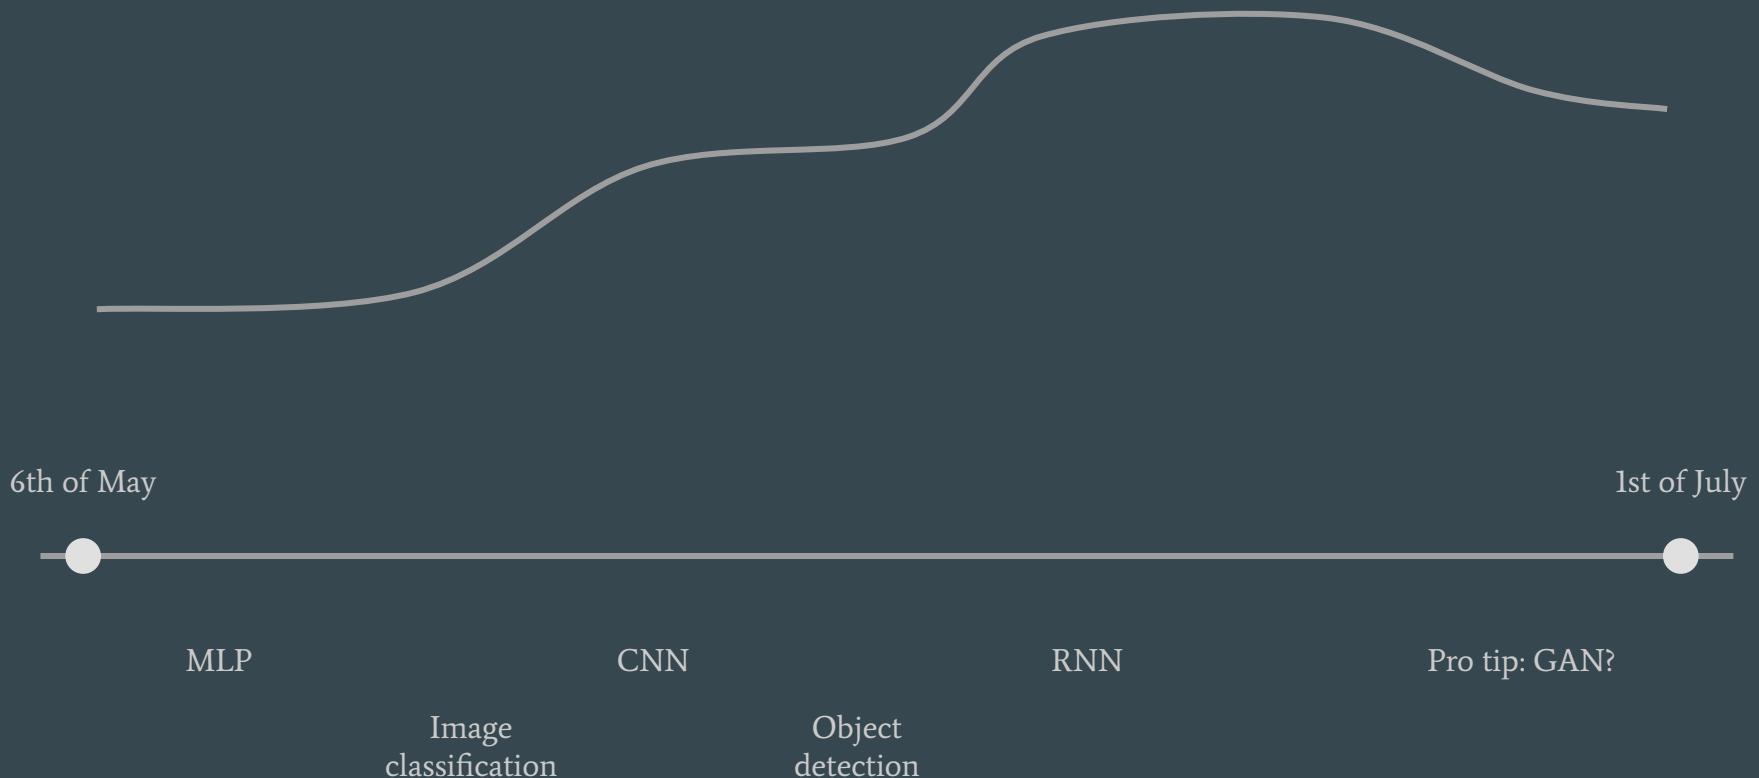
- Add a DL implementation into the roadmap
- Acquire a data scarcity mindset
- It's not all about DL!

# Course overview

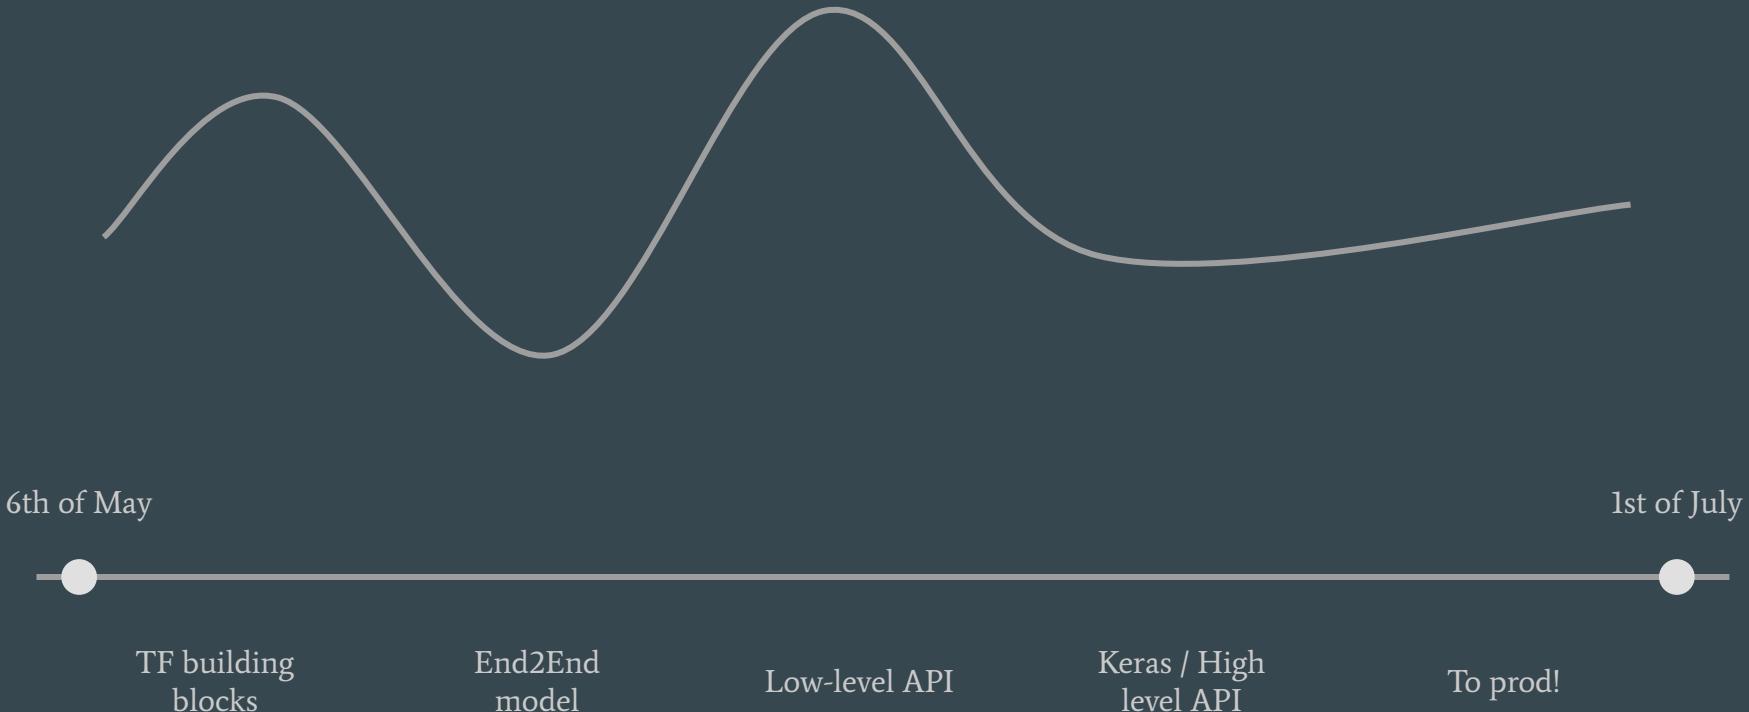
# How is this all going to work?!?!



# Deep learning



# Tensorflow



# Classes

# Classes

What they should be:

- Guide + DIY + review solution → what did we learn?
- Align each one with the subject goals

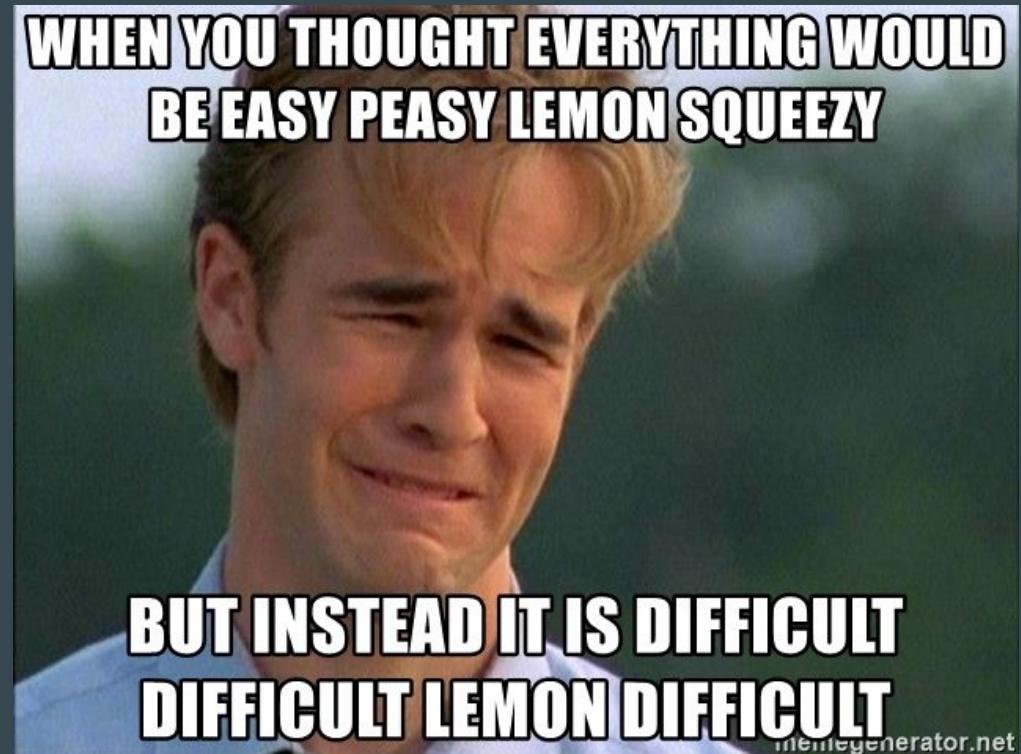
What they should not be:

- Medium blog post
- Tensorflow tutorial
- TF API reference summary

# Classes

Jupyter & Codalab

Old & good python



Begin!

# Computational graph

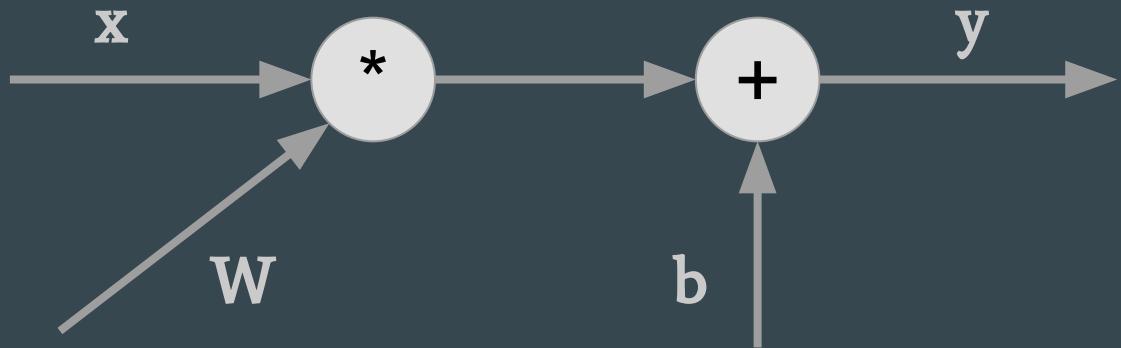
# What is it?

Nodes: operations

Edges: data / tensors

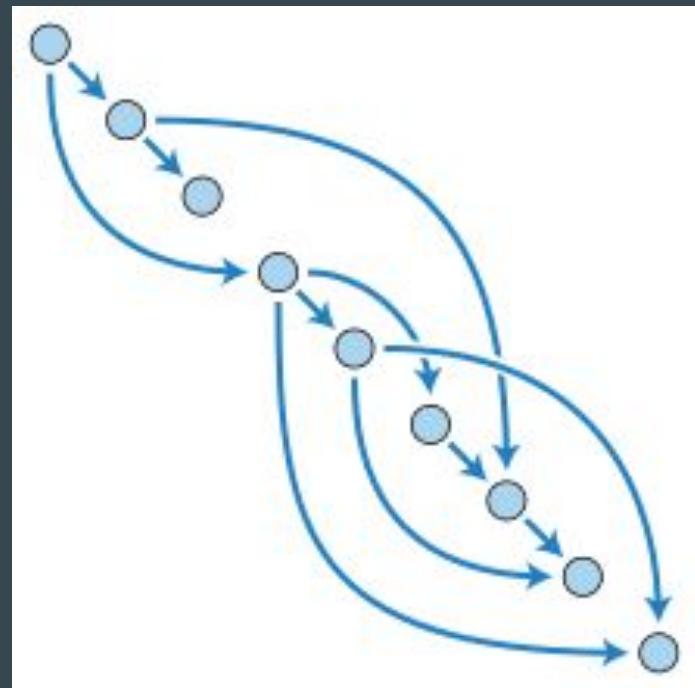
## Example

*Linear regression:*  $y = Wx + b$



# Directed Acyclic Graph (DAG)

- Edged are directed from one node to another
- Nodes have a topological ordering
- No closed loops!



# Static computational graphs

**Define-and-Run methodology.**

**Pros:**

- Speed: can spend a long time optimizing the graph
- Memory: can predict and allocate all mem ahead of time

**Cons:**

- Inflexible: once compiled, it cannot be modified at run time
- Debugging: graph representation doesn't match code
- Static: more difficult to support flexible sized inputs

*Slide credit: Kevin McGuinness (adapted from the original one)*

# Dynamic computational graphs

**Define-by-Run** strategy

**Pros:**

- Can change structure of NN at runtime (add layers, change shape, etc.)
- Support flexible sized inputs
- Easier to debug
- More natural to code

**Cons:**

- Compile at runtime, can be slower
- Dynamic batching more difficult

*Slide credit: Kevin McGuinness (adapted from the original one)*

# Deep learning frameworks



DEEPMLEARNING4J



theano



TensorLayer



TensorFlow

PYTORCH

Caffe



Chainer



Lasagne

mxnet

# Framework basic comparison

Tensorflow:

- Production ready
- Multiple languages supported (Python, C, Java, JavaScript, Go)
- Multi device (Android, iOS, Web browser)
- Mature with a big community
- Tensorboard
- Keras
- Static graph

... Tensorflow 2.0 coming soon! → dynamic graph

Pytorch:

- Pythonic
- Easy to learn
- Framework rather than library
- Integration with Caffe2
- Dynamic graph

... Pytorch 1.0 finally released! → production ready

# Surrounding technologies



# Git

- Course repository (for the coding sessions)
- Use a personal branch to keep track of your progress
- Master branch will be updated each week with the solutions of previous sessions & required material for the current one



# Virtualenv & virtualenvwrapper

- Isolate python environment

# Docker



- Development reproducibility
- Containerize production environment
- TF Serving

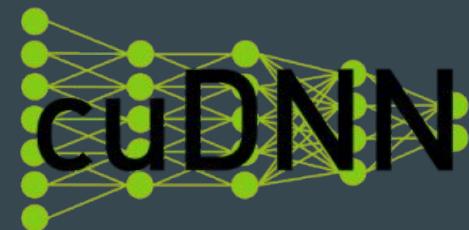


# NVIDIA

**CUDA:** “is a parallel computing platform and application programming interface (API) model created by Nvidia”



**cuDNN:** “is a GPU-accelerated library of primitives for deep neural networks. cuDNN provides highly tuned implementations for standard routines such as forward and backward convolution, pooling, normalization, and activation layers”



# Questions?

